



操作系统原理及应用

李 伟

xchlw@seu.edu.cn

计算机科学与工程学院、软件学院
江苏省网络与信息安全重点实验室



Chapter 3 Processes



Outline

- **Process Concept**
- **Process Scheduling**
- **Operations on Processes**
- **Interprocess Communication**
- **Communication in Client-Server Systems**




Process Concept

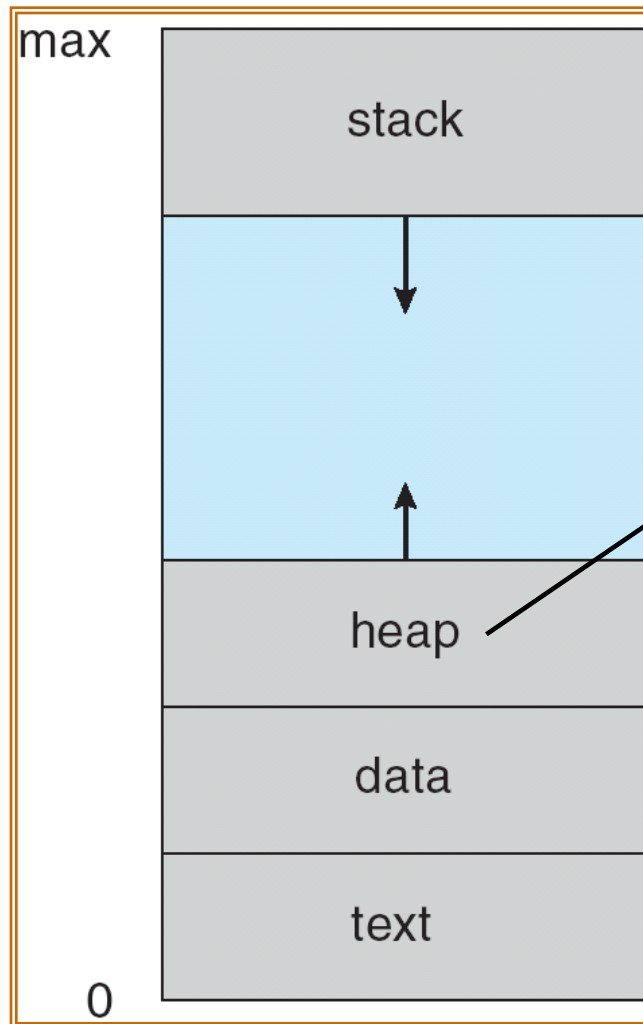
- An operating system executes a variety of programs
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably.



Process Concept

- **Process** – a program in execution; process execution must progress in sequential fashion.
 - **A process includes**
 - text section (program code)
 - **program counter**
 - **contents of the processor's registers**
 - Heap-stack — temporary data (function parameters, return address, local variables)
 - data section — global variables
- current activity
- 

Process in Memory



Heap (堆) is memory that is dynamically allocated during process run time.



Process in Memory

```
//main.cpp
```

```
int a = 0; ← 数据段
```

```
char *p1; ← 数据段
```

```
main()
```

```
{ int b; ← 栈
```

```
char s[] = "abc"; ← 栈
```

```
char *p2; ← 栈
```

```
char *p3 = "123456"; ← 栈
```

```
p1 = (char *)malloc(10); ← 堆
```

```
p2 = (char *)malloc(20); ← 堆
```

```
}
```



Characteristic of Process

- **Dynamic (动态性)**
- **Independency (独立性)**
- **Concurrence (并发性)**
- **Structure (结构化)**



作业1

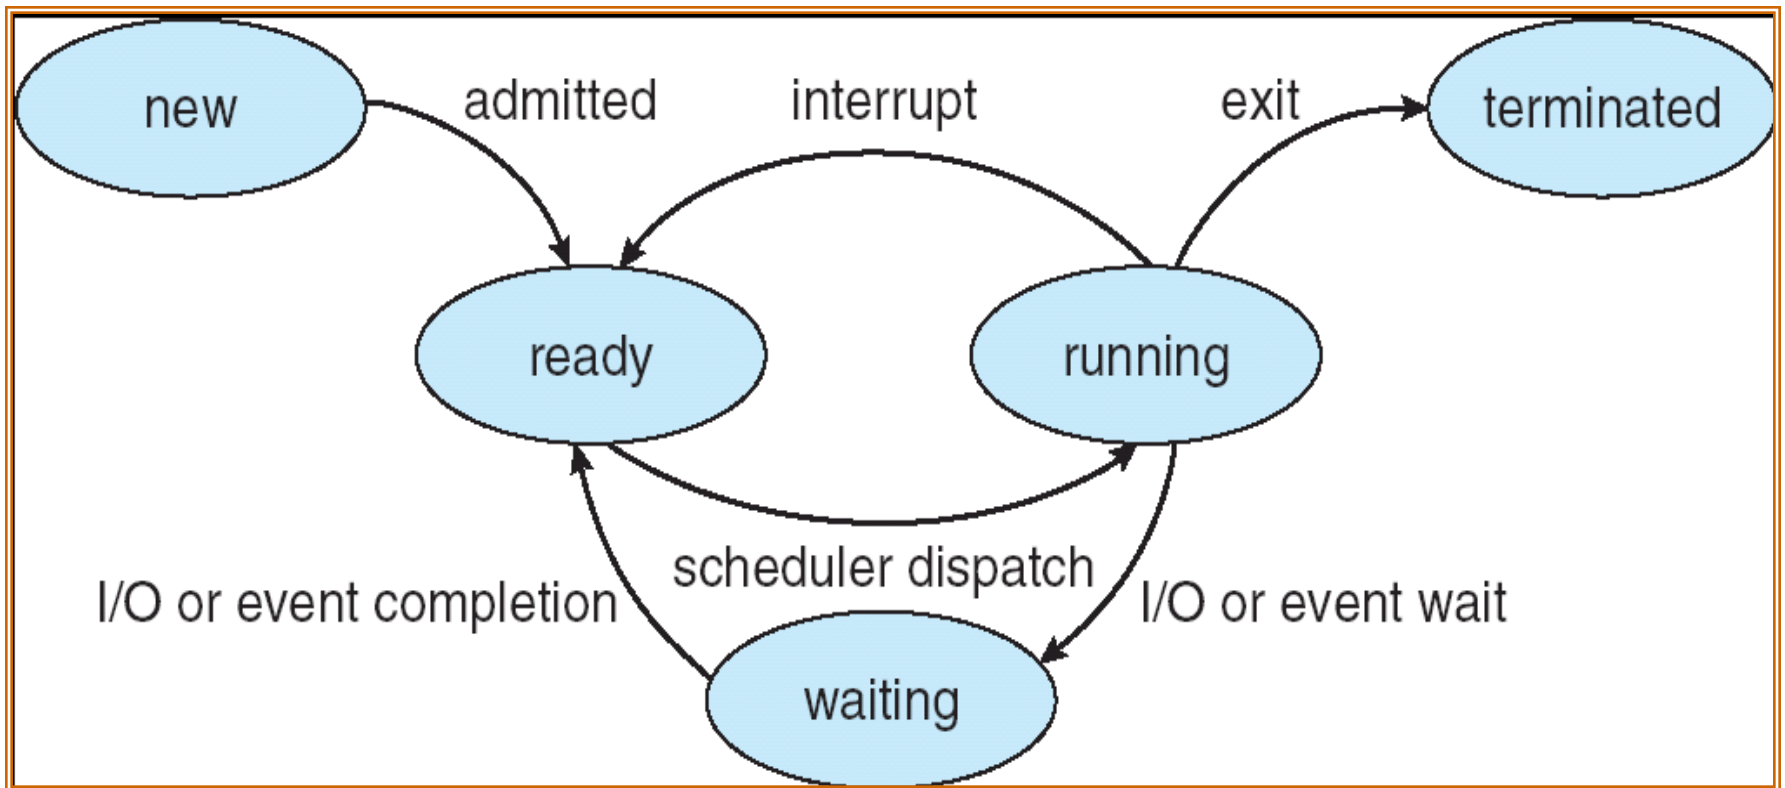
进程和程序是两个密切相关的概念，请阐述它们之间的区别和联系。



Process State

- As a process executes, it changes *state*
 - **New**: The process is being created.
 - **Ready**: The process is waiting to be assigned to a processor.
 - **Running**: Instructions are being executed.
 - **Waiting**: The process is waiting for some event to occur.
 - **Terminated**: The process has finished execution.

Diagram of Process State





Exercise

- 在一个只有单处理机（不考虑多核）的操作系统中，进程有运行、就绪、等待三个基本状态。假如某时刻该系统中有10个用户进程并发执行，且CPU为非核心态情况下，试问：
 - 这时刻系统中处于**运行状态**的用户进程数最多有几个？最少有几个？
 - 这时刻系统中处于**就绪状态**的用户进程数最多有几个？最少有几个？
 - 这时刻系统中处于**等待状态**的用户进程数最多有几个？最少有几个？

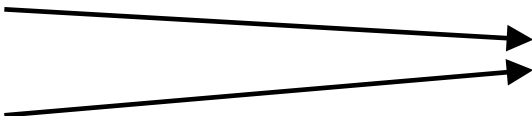


作业2

画出进程在就绪、运行和等待三个基本状态之间的状态转换图，并简述发生相应状态转换的原因。

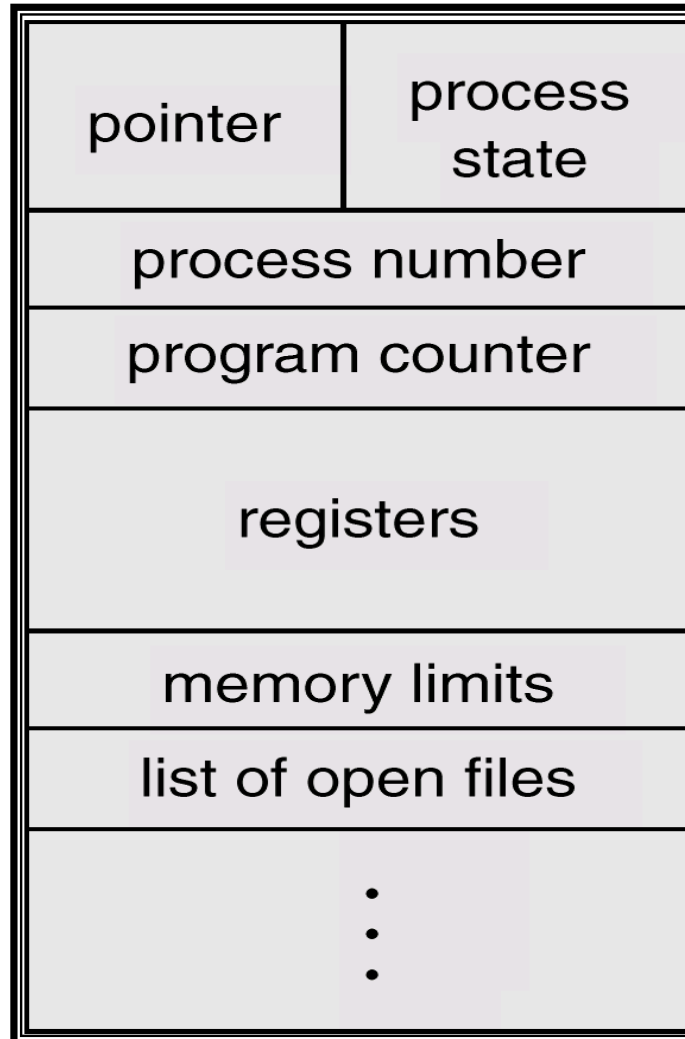


Process Control Block (PCB)

- **Containing the information associated with a specific process**
 - **Process state**
 - **Program counter**
 - **CPU registers**
 - **CPU scheduling information**
 - **Memory-management information**
 - **Accounting information**
 - **I/O status information**
 - **... ..**
- Must be saved when an interrupt occurs**
- 



Process Control Block (PCB)





Outline

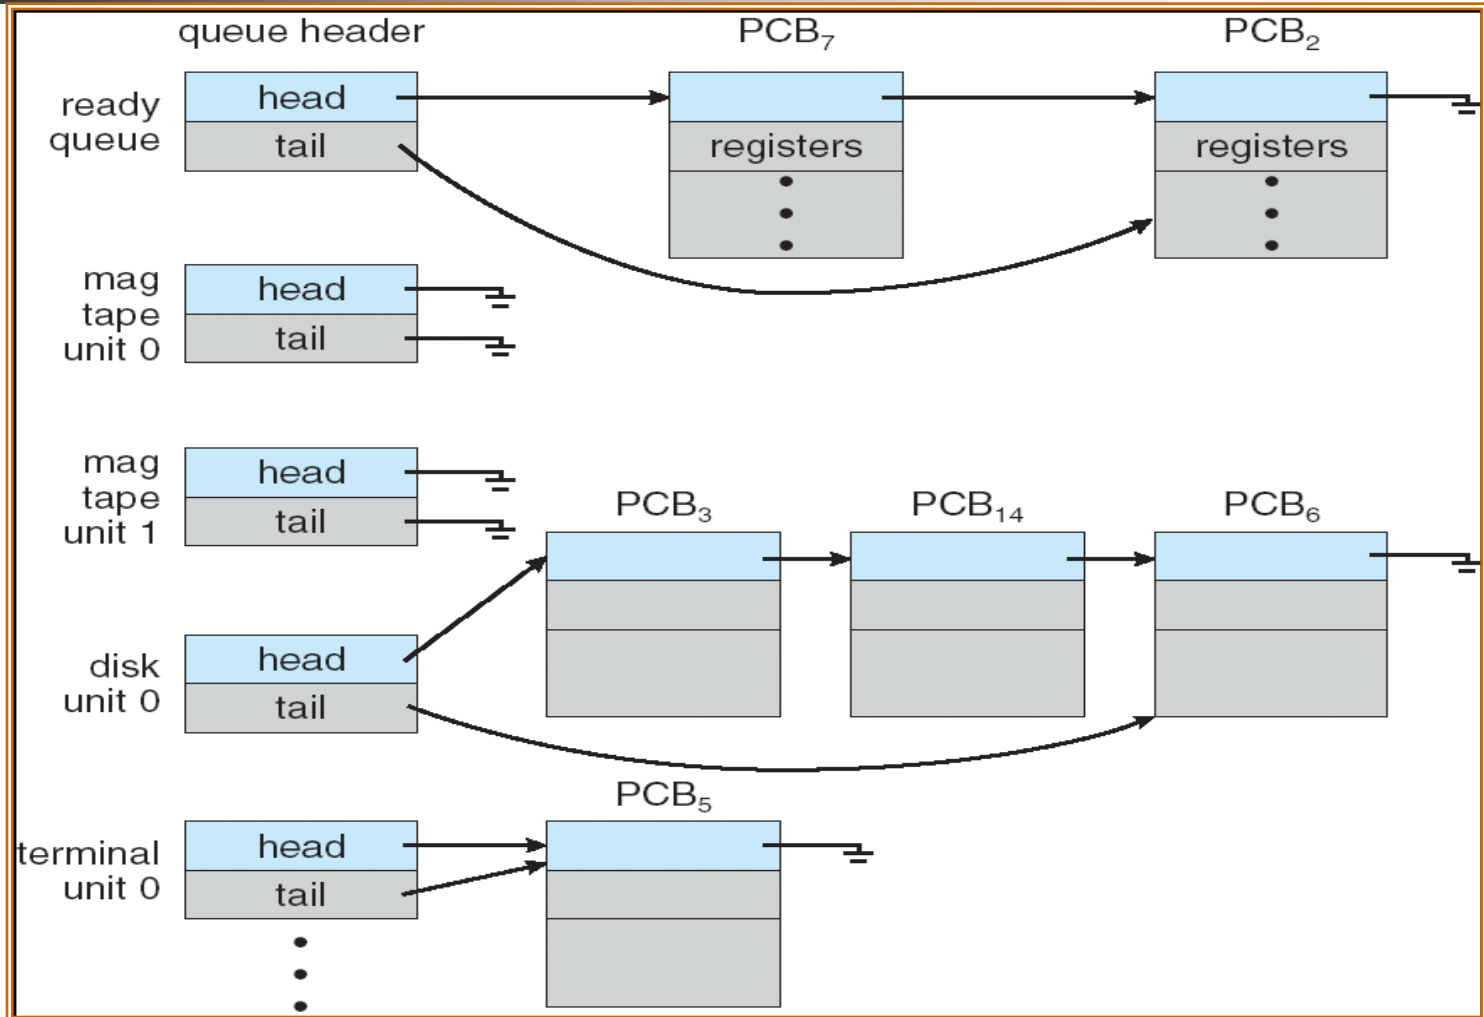
- **Process Concept**
- **Process Scheduling**
- **Operations on Processes**
- **Interprocess Communication**
- **Communication in Client-Server Systems**



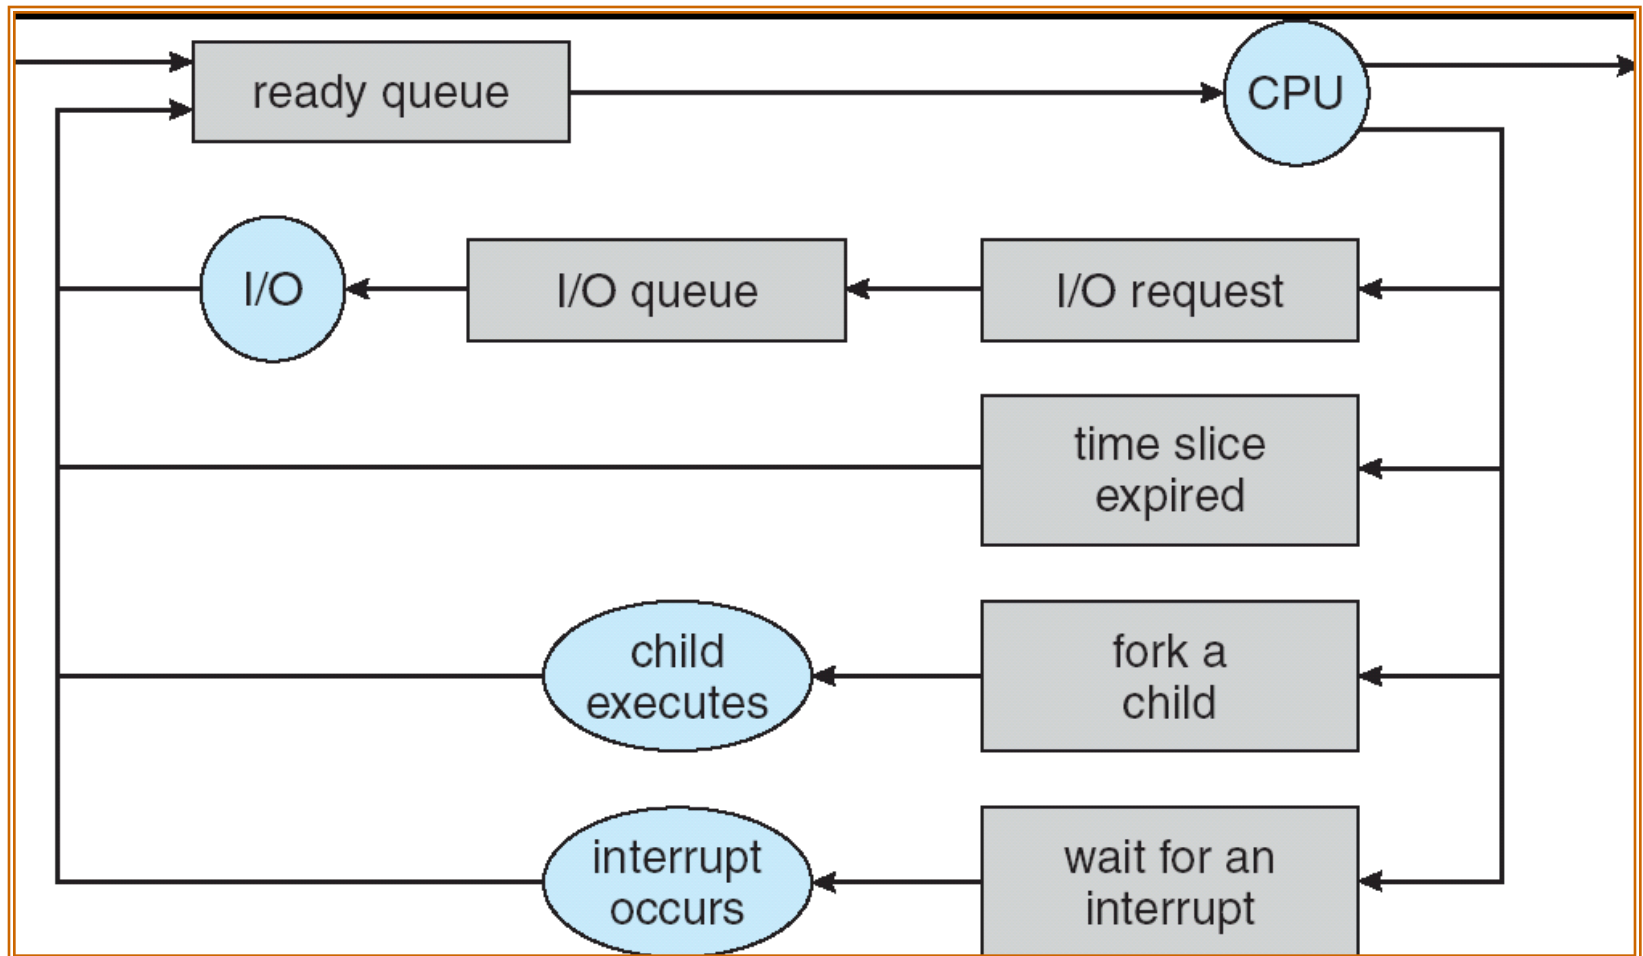
Process Scheduling Queues

- **Job queue** – set of all processes in the system
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** – set of processes waiting for an I/O device
 - Process migration between the various queues

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling





Schedulers

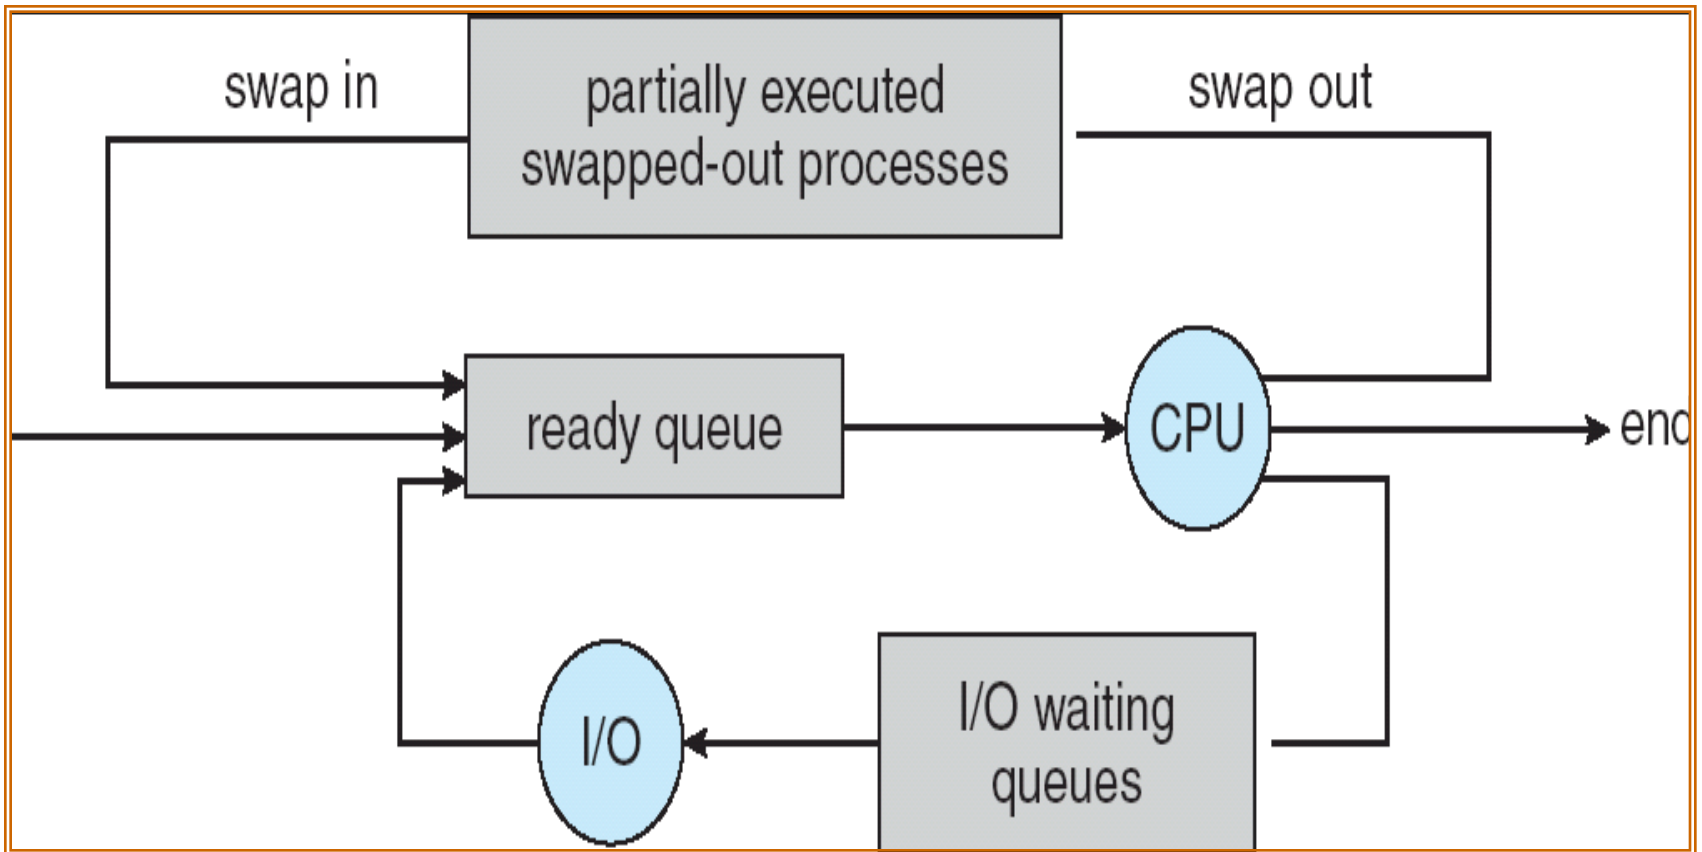
- **Long-term scheduler** (or job scheduler) – selects which processes should be loaded into memory for execution.
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU.
- **Medium-term scheduler** – remove processes from memory and reintroduce them into memory later (**swapping**).



Schedulers

- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow).
- The long-term scheduler controls the **degree of multiprogramming** (*the number of processes in memory*).

Medium-Term Scheduling





Schedulers

- Processes can be described as either
 - ***I/O-bound process*** – spends more time doing I/O than doing computations.
 - ***CPU-bound process*** – spends more time doing computations than doing I/O.
- The long-term scheduler select a good process mix of I/O-bound and CPU-bound processes.



Context Switch

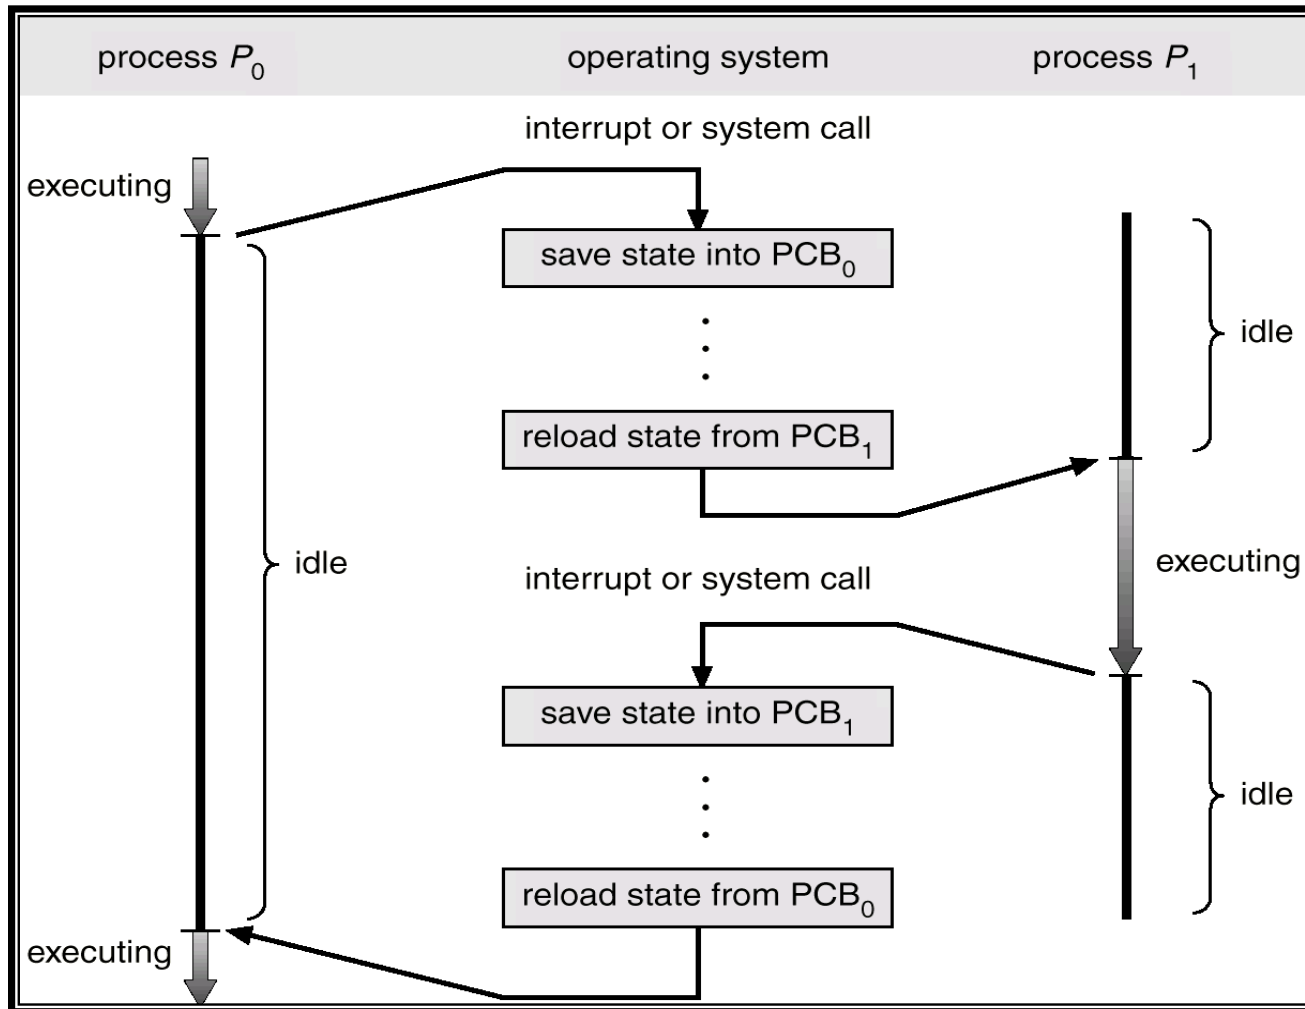
- What is a process context?
 - The **context** of a process is represented in the **PCB** of the process and includes the values of CPU registers, the process state, the program counter, and other memory/file management information.



Context Switch

- **When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.**
- **Context-switch time is overhead; the system does no useful work while switching.**
- **Context-switch time dependent on hardware support.**

CPU Switch From Process to Process





Exercise

- 下列哪一种情况不会引起进程之间的切换？
 - A. 进程调用本程序中定义的函数进行计算
 - B. 进程处理I/O请求
 - C. 进程创建子进程并等待子进程结束
 - D. 产生中断



Outline

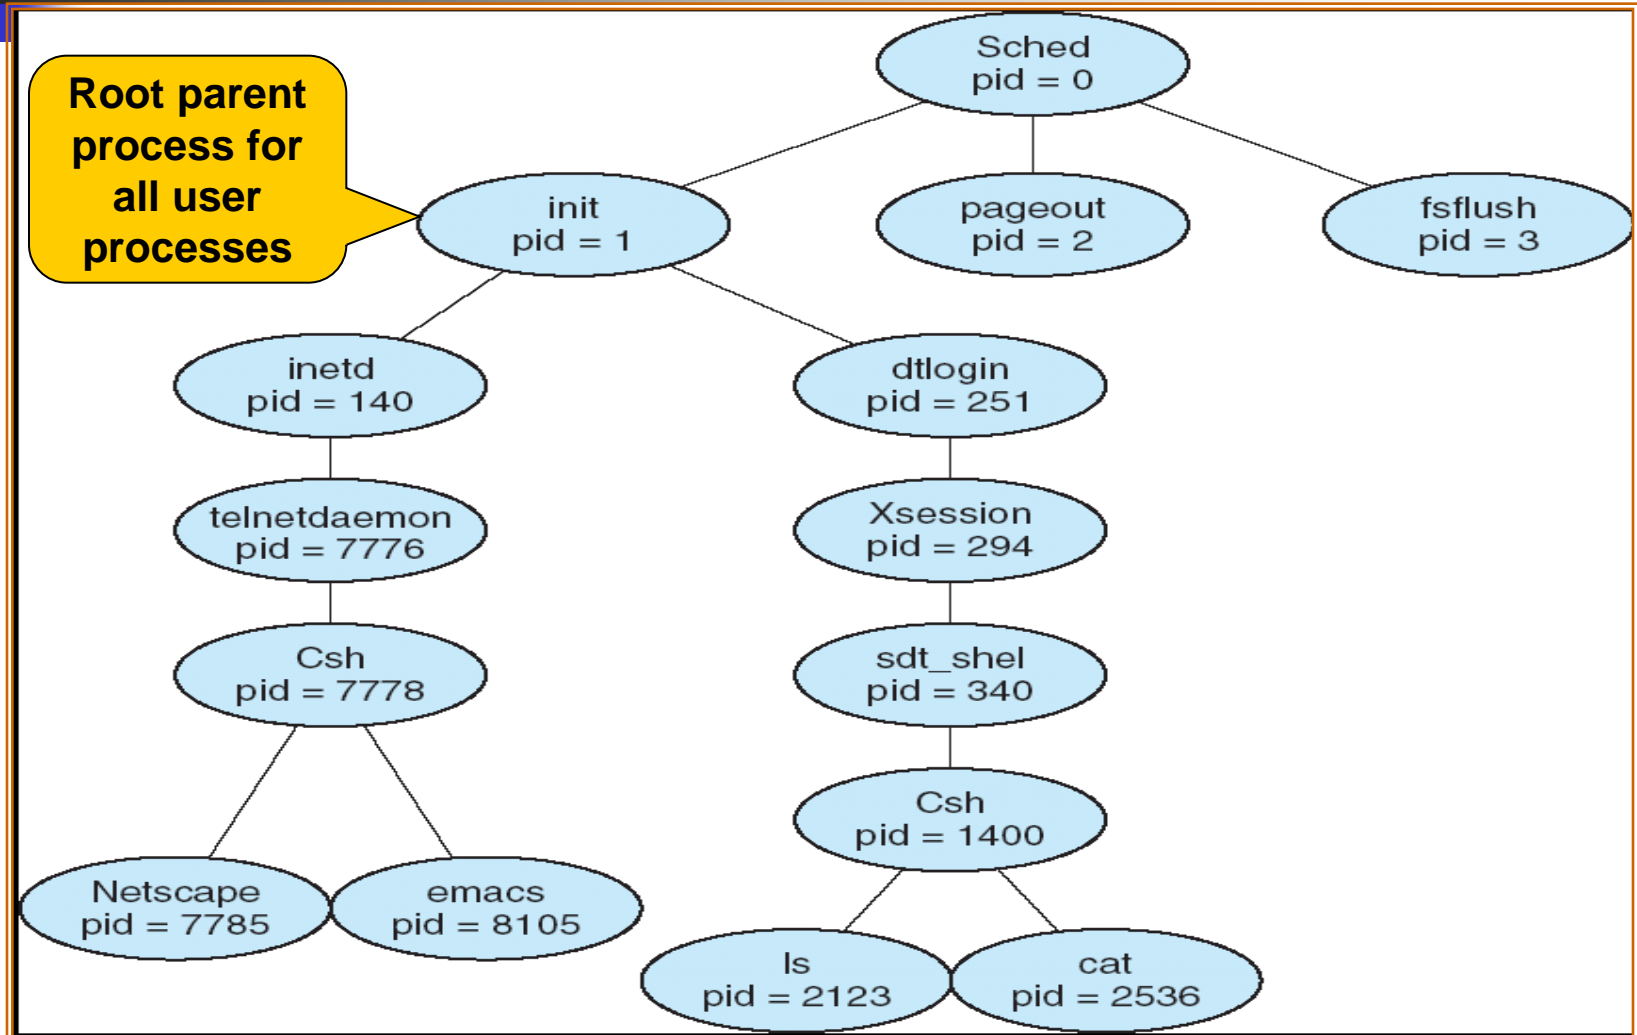
- **Process Concept**
- **Process Scheduling**
- **Operations on Processes**
- **Interprocess Communication**
- **Communication in Client-Server Systems**



Process Creation

- Parent process create children processes, which, in turn create other processes, forming a **tree of processes**.
- **A unique process identifier** (an integer number) is used to identify process.

Processes Tree on Solaris





Process Creation

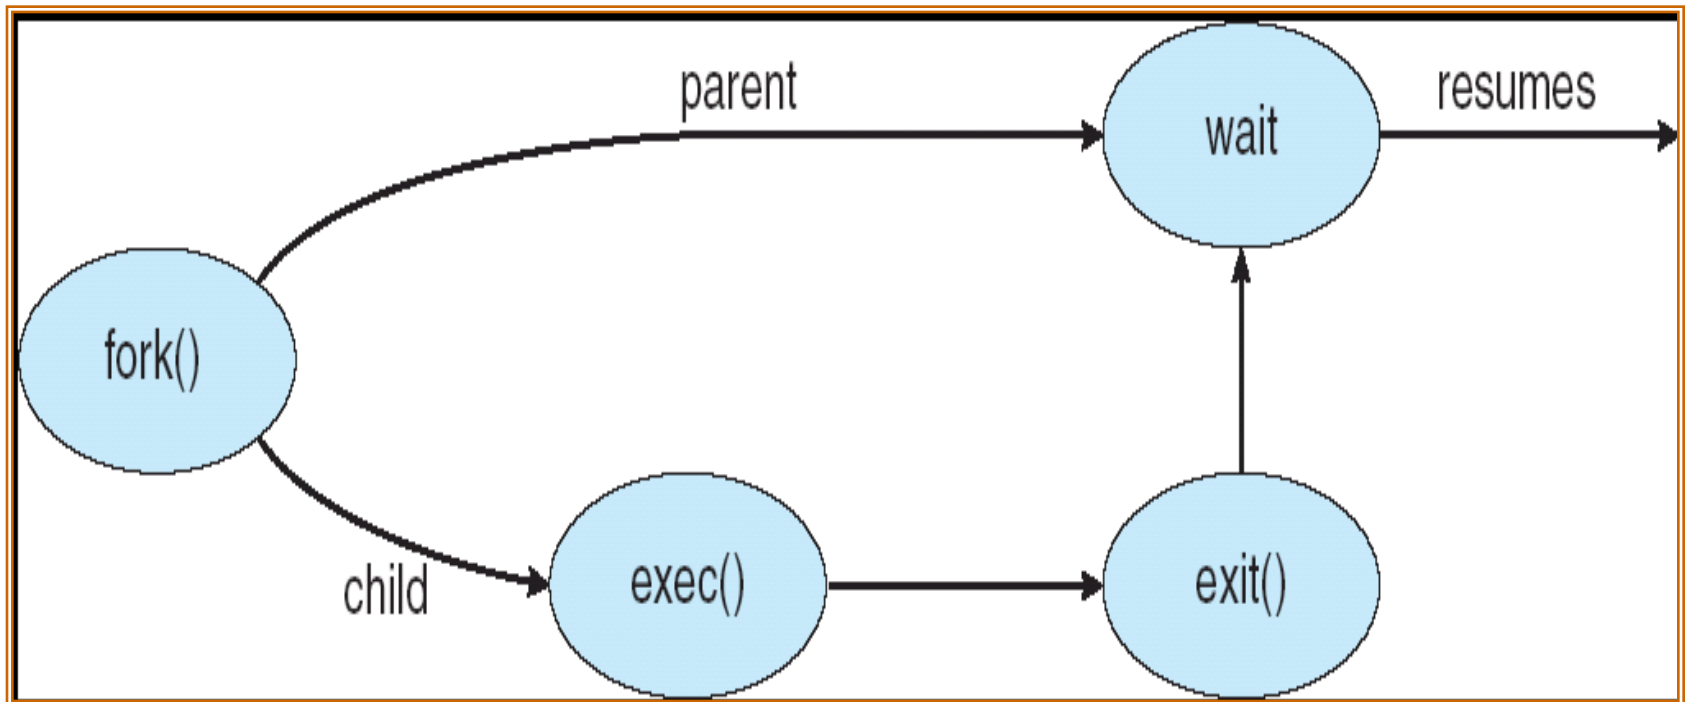
- **Resource sharing**
 - **Parent and children share all resources.**
 - **Children share subset of parent's resources.**
 - **Parent and child share no resources.**
- **Execution**
 - **Parent and children execute concurrently.**
 - **Parent waits until children terminate.**




Process Creation

- **Address space**
 - Child duplicate of parent.
 - Child has a program loaded into it.
- **UNIX examples**
 - **fork() system call** creates new process
 - **exec() system call** used after a fork to replace the process's memory space with a new program.

Process Creation (UNIX)





```
pid = fork();
if (pid < 0){/*error occured*/
    fprintf(stderr, "Fork failed");
    exit(-1);}
else if(pid == 0){/*child process*/
    execlp("/bin/ls", "ls", NULL);
}
else {/*parent process*/
    wait(NULL);
    printf("Child Complete");
    exit(0);
}
```



作业3

■ 下列程序的输出是什么？

```
#include <sys/wait.h>
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
int value = 5;
int main()
{ pid_t pid;
  pid = fork();
  if (pid == 0) { printf("child process, value1 : %d\n", value);
                  value += 15;
                  printf("child process, value2 : %d\n", value);}
  else if (pid > 0) { printf("parent process, value3 : %d\n", value);
                      wait(NULL);
                      printf("parent process, value4 : %d\n", value);
                      exit(0); }
}
```



Process Termination

- **Process executes last statement and asks the operating system to delete it (`exit`).**
 - **Output data from child to parent (via `wait`).**
 - **Process's resources are deallocated by OS.**



Process Termination

- Parent may terminate execution of children processes (**abort**).
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting. (并非所有OS都如此, **UNIX**)
 - Operating system does not allow child to continue if its parent terminates.
 - **Cascading termination (级联终止).**



Outline

- **Process Concept**
- **Process Scheduling**
- **Operations on Processes**
- **Interprocess Communication**
- **Communication in Client-Server Systems**

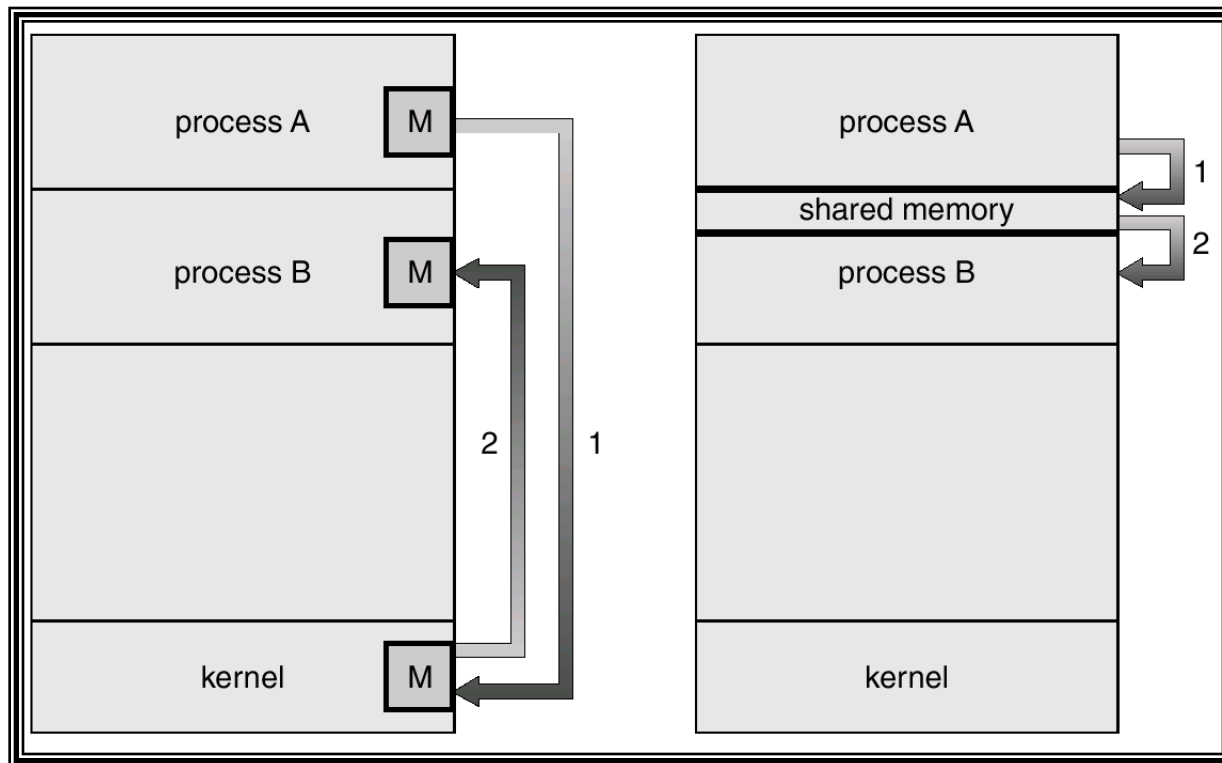


Interprocess Relationship

- ***Independent*** process cannot affect or be affected by the execution of another process.
- ***Cooperating*** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

InterProcess Communication (IPC)

- Communication may take place using either **shared memory** or **message passing**.





Shared-Memory Systems

- **Requiring communicating processes to establish a region of shared memory.**
- **A shared-memory region resides in the address space of the process creating the shared-memory segment.**
- **The processes are responsible for ensuring that they are not writing to the same location simultaneously.**



Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process.
 - *unbounded-buffer* places no practical limit on the size of the buffer.
 - *bounded-buffer* assumes that there is a fixed buffer size.



Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10  
Typedef struct {  
    . . .  
} item;  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;
```



Bounded-Buffer – Producer Process

```
item nextProduced;
```

```
while (1) {
```

```
    while (((in + 1) % BUFFER_SIZE) == out)
```

```
        ; /* do nothing */
```

```
    buffer[in] = nextProduced;
```

```
    in = (in + 1) % BUFFER_SIZE;
```

```
}
```



Bounded-Buffer – Consumer Process

```
item nextConsumed;
```

```
while (1) {
```

```
    while (in == out)
```

```
        ; /* do nothing */
```

```
    nextConsumed = buffer[out];
```

```
    out = (out + 1) % BUFFER_SIZE;
```

```
}
```



Discussion

- Solution is correct, but can only use **`BUFFER_SIZE-1`** elements

Why?

- *如何实现可使用缓冲的最大空间数为 `BUFFER_SIZE`?*



Message-Passing Systems

- MPS facility provides two operations:
 - **send(*message*)** – message size fixed or variable
 - **receive(*message*)**
- If P and Q wish to communicate, they need to:
 - establish a ***communication link*** between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)



Implementation Questions

- **How are links established?**
- **Can a link be associated with more than two processes?**
- **How many links can there be between every pair of communicating processes?**
- **What is the capacity of a link?**
- **Is the size of a message that the link can accommodate fixed or variable?**
- **Is a link unidirectional or bi-directional?**



Direct Communication

- Processes must name each other explicitly:
 - `send(P, message)` – send a message to process *P*
 - `receive(Q, message)` – receive a message from process *Q*
- Properties of communication link
 - Links are established automatically.
 - A link is associated with exactly one pair of communicating processes.
 - Between each pair there exists exactly one link.
 - The link may be unidirectional, but is usually bidirectional.



Indirect Communication

- Messages are sent and received from mailboxes (also referred to as ports).
 - Each mailbox has a unique id.
 - Two processes can communicate only if they share a mailbox.
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes.
 - Each pair of processes may share several communication links.
 - Link may be unidirectional or bidirectional.



Indirect Communication

- **Operations**

- create a new mailbox
- send and receive messages through mailbox
- destroy a mailbox

- **Primitives are defined as:**

send(*A, message*) – send a message to mailbox **A**

receive(*A, message*) – receive a message from mailbox **A**



Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A.
 - P_1 sends; P_2 and P_3 receive.
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



Synchronization

- Message passing may be either blocking or non-blocking.
- **Blocking** is considered synchronous
- **Non-blocking** is considered asynchronous
- send and receive primitives may be either blocking or non-blocking.



Buffering

- Queue of messages attached to the link implemented in one of three ways.
 - **Zero capacity** – 0 messages.
Sender must wait for receiver.
 - **Bounded capacity** – finite length of n messages.
Sender must wait if link full.
 - **Unbounded capacity** – infinite length.
Sender never blocks.



Outline

- **Process Concept**
- **Process Scheduling**
- **Operations on Processes**
- **Interprocess Communication**
- **Communication in Client-Server Systems**



Client-Server Communication

- **Sockets**
- **Remote Procedure Calls**
- **Remote Method Invocation (Java)**

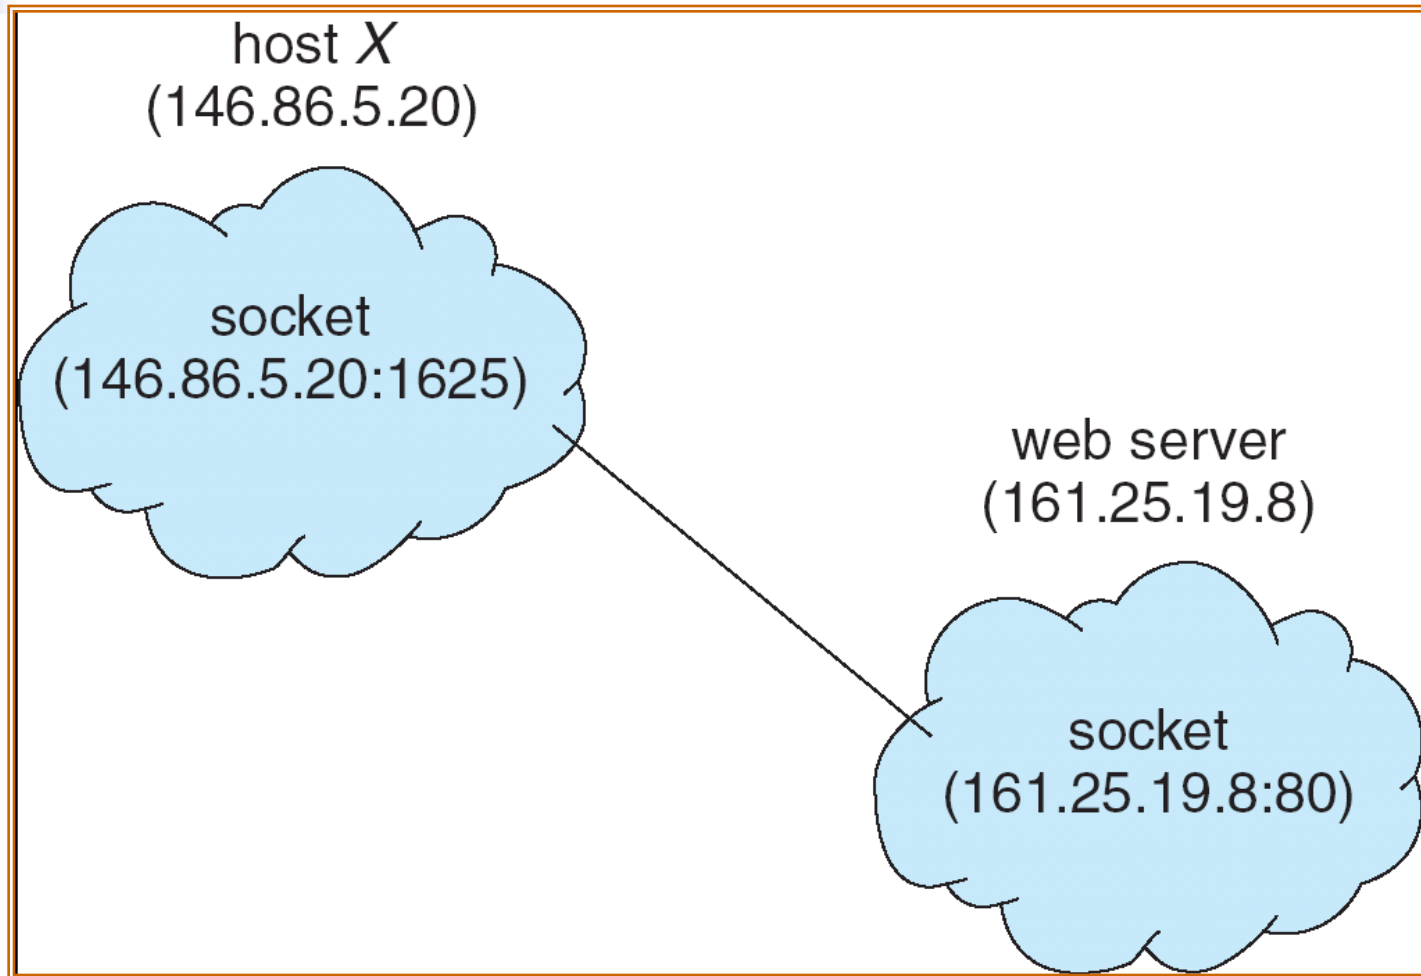


Sockets

- A socket is defined as an *endpoint for communication*.
- Concatenation of IP address and port
- The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- Communication consists between **a pair of sockets**.



Socket Communication





Socket Communication

- All connections must be unique.
- Common and Efficient
- **A low-level form of communication** between distributed processes
- Sockets allow only **an unstructured stream of bytes** to be exchanged between the communicating processes.



Remote Procedure Calls

- **Remote procedure call (RPC)** abstracts procedure calls between processes on networked systems.
- RPC allow a client to invoke a procedure on a remote host as it would invoke a procedure locally.
- **Stubs** – client-side proxy for the actual procedure on the server.
- Each message exchanged contains **an identifier of the function** and **the parameters** to pass to it.



Remote Procedure Calls

- **The client-side stub** locates the port on the server and *marshals* the parameters.
- **The server-side stub** receives this message, unpacks the marshaled parameters, and performs the function on the server.
- Any output is send back to the requester in a separate message.
- Messages exchanged in RPC communication are **well structured**.

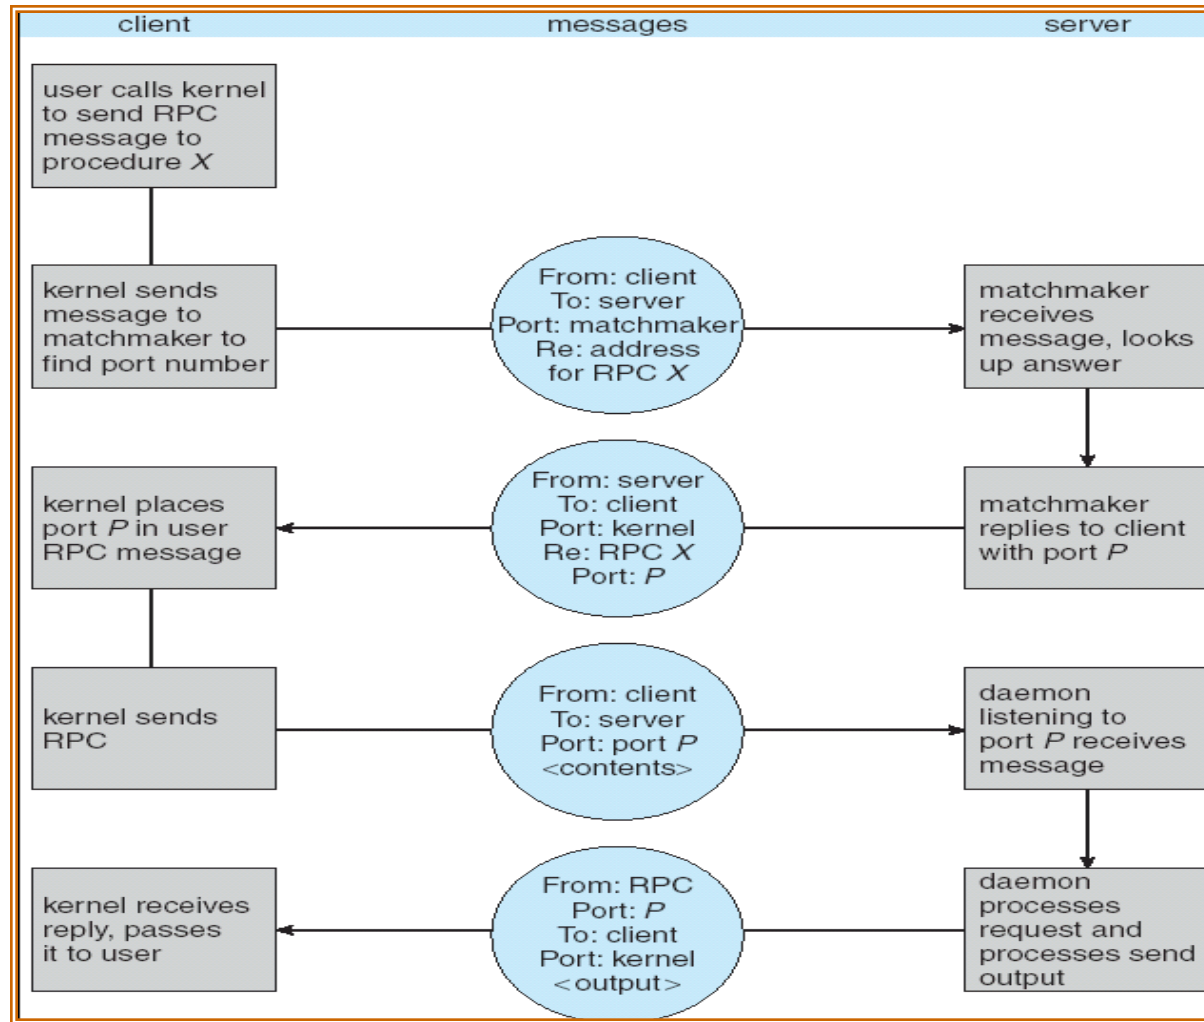


Remote Procedure Calls

- **Some Issues**

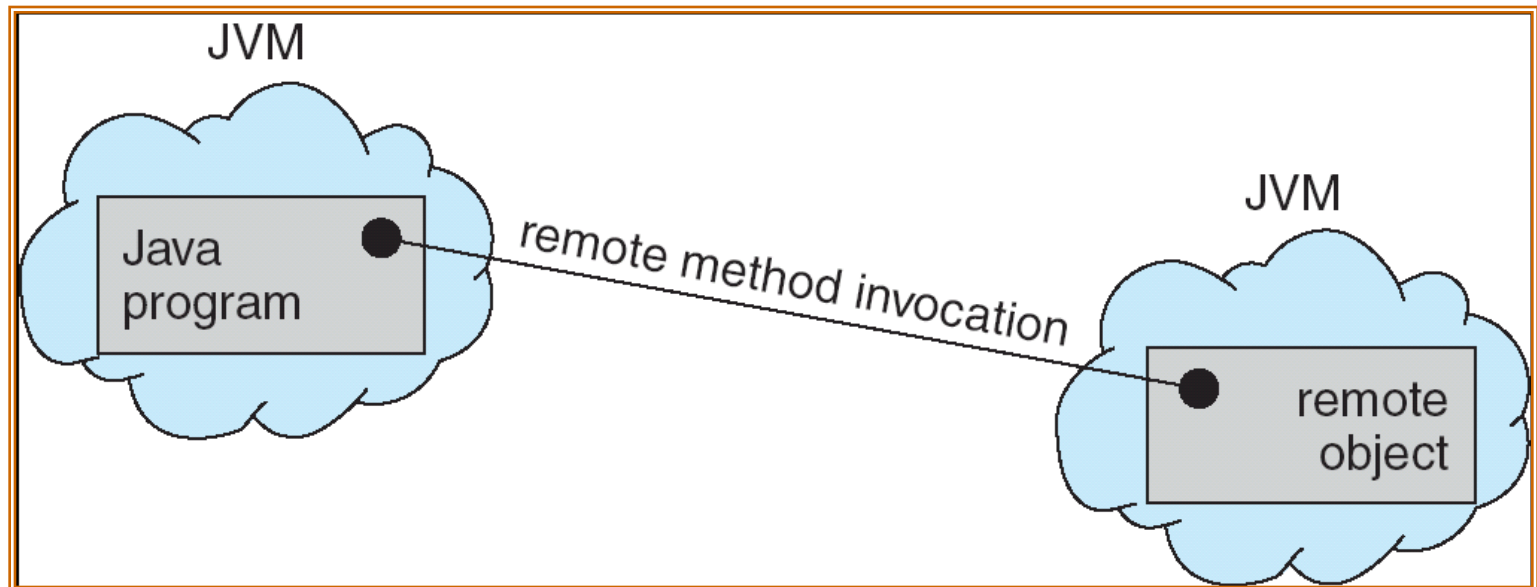
- **The differences in data representation on the client and server machines.**
 - **eXternal Data Representation (XDR)**
- **The semantics of a call**
 - **Exactly once (lost?) : ACK Messages**
 - **At most once (repeated?) : attaching a timestamp to each message.**

Execution of RPC



Remote Method Invocation

- **Remote Method Invocation (RMI)** is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.





Remote Method Invocation

- Differences between RPC and RMI
 - RPC only calls remote **procedures or functions**, and RMI supports invocation of **methods on remote objects**
 - The parameters to remote procedures are **ordinary data structures** in RPC, and that are **objects** in RMI



Remote Method Invocation

- RMI implements the remote object using **stubs** and **skeletons**.
 - Stub is a proxy for the remote object and resides with the client.
 - Skeleton resides with the server, receives the messages from stub and invokes the desired method on the server.

Marshalling Parameters

