



操作系统原理及应用

李 伟

xchlw@seu.edu.cn

计算机科学与工程学院、软件学院
江苏省网络与信息安全重点实验室



Chapter 6 Process Synchronization



Outline

- **Background**
- **The Critical-Section Problem**
- **Synchronization Hardware**
- **Semaphores**
- **Classical Problems of Synchronization**
- **Monitors**
- **Synchronization Examples**



Background

- **Concurrent access to shared data may result in data inconsistency.**
- **Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.**



Background

- Shared-memory solution to **bounded-buffer problem** (Chapter 3) allows at most $n - 1$ items in buffer at the same time.
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
 - We can do so by having an integer **count** that keeps track of the number of full buffers



Bounded-Buffer

- **Shared data**

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```



Bounded-Buffer

- **Producer process**

```
item nextProduced;
```

```
while (1) {
```

```
    while (counter == BUFFER_SIZE);
```

```
        /* do nothing */
```

```
    buffer[in] = nextProduced;
```

```
    in = (in + 1) % BUFFER_SIZE;
```

```
    counter++;
```

```
}
```



Bounded-Buffer

- **Consumer process**

```
item nextConsumed;  
while (1) {  
    while (counter == 0);  
    /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter- -;  
}
```




Bounded-Buffer

- The statements
 counter++;
 counter--;
must be performed ***atomically.***
- Atomic operation means an operation that completes in its entirety without interruption.



Bounded-Buffer

- The statement “**count++**” may be implemented in machine language as:
 - $\text{register}_1 = \text{counter}$
 - $\text{register}_1 = \text{register}_1 + 1$
 - $\text{counter} = \text{register}_1$
- The statement “**count--**” may be implemented as:
 - $\text{register}_2 = \text{counter}$
 - $\text{register}_2 = \text{register}_2 - 1$
 - $\text{counter} = \text{register}_2$



Bounded-Buffer

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.



Bounded-Buffer

- Assume counter is initially 5, what is it finally?

Producer

register1 = counter

register1 = register1 + 1

counter = register1

Consumer

register2 = counter

register2 = register2 - 1

counter = register2



Race Condition

- **Race condition** occurs, if:
 - **two or more** processes/threads access and manipulate the **same** data **concurrently**
 - the outcome of the execution **depends on the particular order** in which the access takes place.
- To prevent race conditions, concurrent processes must be **synchronized**.



Outline

- Background
- **The Critical-Section Problem**
- Synchronization Hardware
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples



The Critical-Section Problem

- n processes all competing to use some shared data
- Each process has **a code segment**, called ***critical section***, in which the shared data is changed.
- **Problem** – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.



The Critical-Section Problem

- Thus, the execution of critical sections must be *mutually exclusive* (e.g., at most one process can be in its critical section at any time).
- The *critical-section problem* is to design a protocol that processes can use to cooperate.

The Critical Section Protocol

do {

⋮

entry section

critical section

exit section

remainder section

⋮

} while (1);

- A critical section protocol consists of **two** parts: an **entry section** and an **exit section**.
- Between them is the critical section that must run in a **mutually exclusive** way.



Solution to Critical-Section Problem

- Any solution to the critical section problem must satisfy the following three conditions:
 - **Mutual Exclusion** (互斥、忙则等待)
 - **Progress** (空闲让进)
 - **Bounded Waiting** (有限等待)
- Moreover, the solution cannot depend on **relative speed** of processes and **scheduling policy**.



Mutual Exclusion

- If a process **P** is executing in its critical section, then **no** other processes can be executing in their critical sections.
- The **critical section protocol** should be capable of blocking processes that wish to enter but cannot.
- Moreover, when the process that is executing in its critical section exits, the **critical section protocol** must be able to know this fact and allows a waiting process to enter.



Progress

- If **no** process is executing in its critical section and some processes wish to enter their critical sections, then
 - Only those processes that are waiting to enter can participate in the competition (to enter their critical sections).
 - No other process can influence this decision.
 - This decision cannot be postponed indefinitely.



Bounded Waiting

- **After** a process made a request to enter its critical section and **before** it is granted the permission to enter, there exists a **bound** on the **number of times** that other processes are allowed to enter.
- Hence, even though a process may be blocked by other waiting processes, **it will not be waiting forever.**
- **Assume that each process executes at a nonzero speed**
- **No assumption concerning relative speed of the n processes**



Initial Attempts to Solve Problem

- Only 2 processes, P_0 and P_1
- General structure of process P_i (other process P_j)
do {
 entry section
 critical section
 exit section
 remainder section
} while (1);
- Processes may share some common variables to synchronize their actions.



Algorithm 1

- Shared variables:
 - **int turn;**
initially **turn = i (or turn=j)**
- Process P_i :
 - do { while (turn != i) ;**
critical section
turn = j;
remainder section
} while (1);
- are forced to run in an alternating way.
- Satisfies **mutual exclusion**, but not **progress**



Algorithm 2

- Shared variables
 - **boolean** **flag[2];**
initially **flag[0] = flag[1] = false.**
 - **flag [i] = true** $\Rightarrow P_i$ ready to enter its critical section
- Process P_i
 - do {
 - flag[i] = true;**
 - while (flag[j]) ;**
 - critical section
 - flag [i] = false;**
 - remainder section
 - } while (1);**
- Satisfies **mutual exclusion**, but not **progress**.



Is the following algorithm correct?

- Shared variables

- `boolean flag[2];`

- initially `flag[0] = flag[1] = false.`

- Process P_i :

- ```
do {
 while (flag[j]) ;
 flag[i] = true;
 critical section
 flag [i] = false;
 remainder section
} while (1);
```



## Algorithm 3—Peterson's Solution

---

- Combined shared variables of algorithms 1, 2.
- Process  $P_i$ 
  - do {
    - $\text{flag}[i] = \text{true};$
    - $\text{turn} = i;$
    - $\text{while} (\text{flag}[j] \text{ and } \text{turn} == j) ;$
    - critical section
    - $\text{flag}[i] = \text{false};$
    - remainder section
  - } while (1);
- Meets all three requirements; solves the critical-section problem **for two processes.**



# Bakery Algorithm

---

## Critical section for n processes

- Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- If processes  $P_i$  and  $P_j$  receive the same number, if  $i < j$ , then  $P_i$  is served first; else  $P_j$  is served first.
- The numbering scheme always generates numbers in non-decreasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...



# Bakery Algorithm

---

- **Notation**

- $(a,b) < (c,d)$  if  $a < c$  or if  $a = c$  and  $b < d$
- $\max(a_0, \dots, a_{n-1})$  is a number  $k$ , such that  $k \geq a_i$  for  $i$  from 0 to  $n - 1$

- **Shared data**

**boolean choosing[n];**

**int number[n];**

**Data structures are initialized to false and 0 respectively**



# Bakery Algorithm

---

do {

    choosing[ i ] = true;

    number[ i ] = max(number[0], number[1], ..., number [n – 1])+1;

    choosing[ i ] = false;

    for (j = 0; j < n; j++) {

        while (choosing[ j ]) ;

        while ((number[ j ] != 0) && ((number[ j ],j) < (number[ i ],i))) ;

    }

    critical section

    number[ i ] = 0;

    remainder section

} while (1);

*Discussion*



# Outline

---

- Background
- The Critical-Section Problem
- **Synchronization Hardware**
- Semaphores
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples



# Hardware Support

---

- **There are two types of hardware synchronization supports**
  - **Disabling/Enabling interrupts**
    - This is slow and difficult to implement on multiprocessor systems.
  - **Special machine instructions**
    - TestAndSet (TS)
    - Swap

# Interrupt Disabling

```
do {
```



entry

disable interrupts

critical section

enable interrupts



exit

```
} while (1);
```

- Because interrupts are disabled, no context switch will occur in a critical section.
- Infeasible in a multiprocessor system because all CPUs must be informed.
- Some features that depend on interrupts (e.g., clock) may not work properly.





# TestAndSet

---

- **Test** and **modify** the content of a word **atomically**

```
boolean TestAndSet(boolean &target)
{
 boolean rv = ⌖
 &target = true;
 return rv;
}
```



# Mutual Exclusion with TestAndSet

---

- Shared data:  
    boolean lock = false;
- Process  $P_i$   
    do {  
        while (TestAndSet(lock)) ;  
        critical section  
        lock = false;  
        remainder section  
    }



# Swap

---

- **Atomically** swap two variables.

```
void Swap(boolean &a, boolean &b)
```

```
{
```

```
 boolean temp = &a;
```

```
 &a = &b;
```

```
 &b = temp;
```

```
}
```



# Mutual Exclusion with Swap

---

- **Global** Shared data  
boolean lock; (initialized to false):
- **Local** variable for each process  
boolean key;
- Process  $P_i$   
do {  
    key = true;  
    while (key == true)  
        Swap(lock, key);  
    critical section  
    lock = false;  
    remainder section  
}



# Satisfying Three Conditions with TestAndSet

Initially Boolean `waiting[i] = false`; `lock=false`

## Entry Section

```
waiting[i] = true;
key = true;
while (waiting[i] && key)
 key = TestAndSet(lock);
waiting[i] = false;
```

如果删除该语句，  
会产生什么后果？

## Exit Section

```
j = (i+1)%n
while ((j!=i) && !waiting[j])
 j = (j+1)%n;
if (j == i)
 lock = false;
else
 waiting[j] = false;
```



# Outline

---

- Background
- The Critical-Section Problem
- Synchronization Hardware
- **Semaphores**
- Classical Problems of Synchronization
- Monitors
- Synchronization Examples



# Semaphores

---

- Synchronization tool
- Semaphore  $S$  – integer variable
- can only be accessed via two standard **atomic** operations

*wait* ( $S$ ):  $P()$

while  $S \leq 0$ ;  
     $S--$ ;

*signal* ( $S$ ):  $V()$

$S++$ ;



# Two Types of Semaphores

---

- ***Counting semaphore*** – integer value can range over an unrestricted domain.
- ***Binary semaphore*** – integer value can range only between 0 and 1; also called **mutex locks**.
- Can implement a counting semaphore  $S$  as a binary semaphore.





# Critical Section of n Processes

---

- Shared data:

semaphore mutex; //initially *mutex* = 1

- Process  $P_i$ :

do {

    wait(mutex);

    critical section

    signal(mutex);

    remainder section

} while (1);



**Wrong or Right?**



# Semaphore as a General Synchronization Tool

---

- Execute  $B$  in  $P_j$  only after  $A$  executed in  $P_i$
- Use semaphore *flag* initialized to 0
- Code:

|                      |                    |
|----------------------|--------------------|
| $P_i$                | $P_j$              |
| $\vdots$             | $\vdots$           |
| $A$                  | <i>wait (flag)</i> |
| <i>signal (flag)</i> | $B$                |



# Semaphore Implementation

---

- The main disadvantage of the above classical semaphore definition is that it requires **busy waiting**. (while a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the entry section)
- This type of semaphore is also called a **spinlock** (自旋锁) because the process spins while waiting for the lock.
- To overcome the need for busy waiting, the process can **block itself** rather than engaging in busy waiting.



# Semaphore Implementation

---

- Define a semaphore as a record

```
typedef struct {
 int value;
 struct process *L; // waiting queue
} semaphore;
```
- Assume two simple operations:
  - **block()** suspends the process that invokes it.
  - **wakeup(*P*)** resumes the execution of a blocked process *P*.



# Implementation

---

- Semaphore operations now defined as

*wait(S):*

S.value--;

if (S.value < 0) {

add this process to S.L;

block();

}

*signal(S):*

S.value++;

if (S.value <= 0) {

remove a process P from S.L;

wakeup(P);

}



# Implementation

---

- If the semaphore value is negative, its magnitude is the number of processes waiting on that semaphore.
- The critical aspect of semaphores is that they be executed **atomically**. (**Critical-section Problem**)
- Busy waiting has **not** been **completely eliminated**.
  - Busy waiting has been removed from the critical sections of application programs.
  - Furthermore, we have limited busy waiting to the critical sections of the **wait()** and **signal()** operations.



# Exercises

---

- 5个进程共享某一临界资源，则互斥信号量的取值范围是多少？
- 有4个进程共享一程序段，而每次最多允许2个进程进入该程序段，则信号量的初值是多少？



# Deadlock and Starvation

---

- **Deadlock** – **two or more** processes are waiting indefinitely for **an event** that can be caused by only one of the waiting processes.
- Let S and Q be two semaphores all initialized to 1

|                   |                   |
|-------------------|-------------------|
| $P_0$             | $P_1$             |
| <i>wait(S);</i>   | <i>wait(Q);</i>   |
| <i>wait(Q);</i>   | <i>wait(S);</i>   |
| $\vdots$          | $\vdots$          |
| <i>signal(S);</i> | <i>signal(Q);</i> |
| <i>signal(Q)</i>  | <i>signal(S);</i> |

- **Starvation** – indefinite blocking. **A process** may never be removed from the semaphore queue in which it is suspended.





# Outline

---

- Background
- The Critical-Section Problem
- Synchronization Hardware
- Semaphores
- **Classical Problems of Synchronization**
- Monitors
- Synchronization Examples



# Classical Problems of Synchronization

---

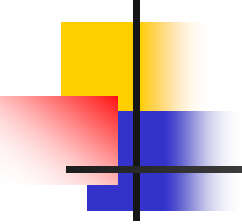
- **Bounded-Buffer Problem**
- **Readers and Writers Problem**
- **Dining-Philosophers Problem**



# Bounded-Buffer Problem

---

- **Shared data**
  - **Semaphore: full, empty, mutex;**
  - **Initially: full = 0, empty = n, mutex = 1**



# Bounded-Buffer Problem

## Producer Process

---

```
do { ...
 produce an item in nextp
 ...
 wait(empty);
 wait(mutex);
 ...
 add nextp to buffer
 ...
 signal(mutex);
 signal(full);
} while (1);
```



# Bounded-Buffer Problem

## Consumer Process

---

```
do {
 wait(full)
 wait(mutex);
 ...
 remove an item from buffer to nextc
 ...
 signal(mutex);
 signal(empty);
 ...
 consume the item in nextc
 ...
} while (1);
```



# 作业1

---

- 在生产者——消费者问题中，信号量 `mutex`, `empty`, `full` 的作用是什么？如果分别对调生产者进程中的两个 `wait` 操作和两个 `signal` 操作，则可能发生什么情况？



# Readers-Writers Problem

---

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write.
- **Problem** – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
  - First readers-writers problem——**Reader first**
  - Second readers-writers problem——**Writer first**



# First Readers-Writers Problem

---

- **Shared data**

**int readcount;**

**semaphore wrt, mutex;**

- **Initially**

**mutex = 1, wrt = 1, readcount = 0**





# First Readers-Writers Problem

## Writer Process

---

```
do {
 wait(wrt);
 ...
 writing is performed
 ...
 signal(wrt);
} while (1)
```



# First Readers-Writers Problem

## Reader Process

---

```
do {
 wait(mutex);
 readcount++;
 if (readcount == 1)
 wait(wrt);
 signal(mutex);
 ...
 reading is performed
 ...
 wait(mutex);
 readcount--;
 if (readcount == 0)
 signal(wrt);
 signal(mutex);
}
```



**Wrong or Right?**




## 作业2

---

- 用信号量解决无饥饿的读者——写者问题。

# Dining-Philosophers Problem



- **Shared data**  
**semaphore chopstick[5];**  
**Initially all values are 1**



# Dining-Philosophers Problem

---

- Philosopher  $i$ :

```
do {
 wait(chopstick[i])
 wait(chopstick[(i+1) % 5])

 ...
 eat

 ...
 signal(chopstick[i]);
 signal(chopstick[(i+1) % 5]);

 ...
 think

 ...
} while (1);
```



**Wrong or Right?**



# Exercise

---

- 用信号量解决“独木桥”问题：同一方向的行人可连续过桥，当某一方向有人过桥时，另一方向的人必须等待；当某一方向无人过桥时，另一方向的行人可以过桥。



# Exercise

---

## ■ Shared data

**int countA = 0** //表示A方向上已在独木桥上的行人数目

**int countB = 0** //表示B方向上已在独木桥上的行人数目

**semaphore MA =1** //实现对countA的互斥修改

**MB =1** //实现对countB的互斥修改

**mutex =1** //实现两个方向的行人对独木桥  
的互斥使用



# Exercise

---

A方向过桥进程:

```
do{
 wait(MA);
 countA ++;
 if (countA == 1) then wait(mutex);
 signal(MA);
 过桥;
 wait(MA);
 countA --;
 if (countA == 0) then signal(mutex);
 signal(MA)
} while (1);
```

**B方向过桥进程?**





# Outline

---

- **Background**
- **The Critical-Section Problem**
- **Synchronization Hardware**
- **Semaphores**
- **Classical Problems of Synchronization**
- **Monitors**
- **Synchronization Examples**



# Problems with Semaphores

---

- **Incorrect use of semaphore operations**
  - **signal (mutex) .... wait (mutex)**
  - **wait (mutex) ... wait (mutex)**
  - **Omitting of wait (mutex) or signal (mutex) (or both)**

**把分散在各进程中的临界区集中起来管理**



# Monitors

---

- **High-level synchronization construct** that allows the safe sharing of an abstract data type among concurrent processes.

```
monitor monitor-name
{
 shared variable declarations
 procedure body P1 (...) {
 ...
 }
 procedure body P2 (...) {
 ...
 }
 procedure body Pn (...) {
 ...
 }
 { initialization code }
}
```




# Monitors: Mutual Exclusion

---

- ***No more than one process*** can be executing ***within*** a monitor.
- When a process calls a monitor procedure and the monitor has a process running, the caller will be **blocked outside of the monitor**.
- Thus, ***mutual exclusion*** is guaranteed within a monitor.

# Schematic View of a Monitor






# Condition Variables

---

- To allow a process to wait within the monitor, a **condition** variable must be declared, as **condition x, y;**
- Condition variable can only be used with the operations **wait** and **signal**.
  - **x.wait()** means that the process invoking this operation is suspended until another process invokes **x.signal()**;
  - The **x.signal()** operation resumes exactly one suspended process. If no process is suspended, then the **signal()** operation has no effect.

# Monitor With Condition Variables





# Condition Variables

---

- Consider the released process and the signaling process
  - There are **two** processes executing in the monitor, and mutual exclusion is violated!
- Approaches to address this problem
  - The released process takes the monitor and the signaling process waits somewhere. (唤醒并等待)
  - The released process waits somewhere and the signaling process continues to use the monitor. (唤醒并继续)





# Semaphore vs. Condition

| Semaphores                                                                                           | Condition Variables                                                                                                  |
|------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|
| Can be used anywhere, but not in a monitor                                                           | Can only be used in monitors                                                                                         |
| <code>wait()</code> does not always block its caller                                                 | <code>wait()</code> <b>always</b> blocks its caller                                                                  |
| <code>signal()</code> either releases a process, or increases the semaphore counter                  | <code>signal()</code> either releases a process, or the signal is <b>lost</b> as if it never occurs                  |
| If <code>signal()</code> releases a process, the caller and the released <b><i>both continue</i></b> | If <code>signal()</code> releases a process, either the caller or the released continues, but <b><i>not both</i></b> |



# Monitor vs. Process

---

- 管程定义的是公用数据结构，而进程定义的是私有数据结构
- 管程把共享变量上的同步操作集中起来，而临界区却分散在每个进程中
- 管程是为管理共享资源而建立的，进程主要是为占有系统资源和实现系统并发性而引入的
- 管程是被欲使用共享资源的进程所调用，管程和调用它的进程不能并发工作，而进程之间能并发工作
- 管程是语言或操作系统的成分，不必创建或撤销，而进程有生命周期，由创建而产生至撤销便消亡



# Dining Philosophers

---

**monitor dp**

**{**

**enum {thinking, hungry, eating} state[5];**

**condition self[5];**

**void pickup(int i)**

**// following slides**

**void putdown(int i)**

**// following slides**

**void test(int i)**

**// following slides**

**void init()**

**{ for (int i = 0; i < 5; i++)**

**state[i] = thinking; }**

**}**



# Dining Philosophers

---

```
void pickup(int i)
{
 state[i] = hungry;
 test[i];
 if (state[i] != eating)
 self[i].wait();
}
```

```
void putdown(int i)
{
 state[i] = thinking;
 // test left and right neighbors
 test((i+4) % 5);
 test((i+1) % 5);
}
```



# Dining Philosophers

---

```
void test(int i)
{
 if ((state[(i + 4) % 5] != eating) &&
 (state[i] == hungry) &&
 (state[(i + 1) % 5] != eating))
 {
 state[i] = eating;
 self[i].signal();
 }
}
```



# Dining Philosophers

---

- Each philosopher  $i$  invokes the operations **pickup()** and **putdown()** in the following sequence:

**dp.pickup (i)**

**EAT**

**dp.putdown (i)**

# Monitor Implementation Using Semaphores

- Variables

确保管程的互斥调用

```
semaphore mutex; // (initially = 1)
semaphore next; // (initially = 0)
int next-count = 0;
```

- Each external procedure  $F$  will be replaced by

管程内挂起进程的总数

```
wait (mutex);
```

```
...
```

```
body of F ;
```

```
...
```

```
if (next-count > 0)
```

```
 signal (next)
```

```
else signal (mutex);
```

防止执行  
Signal操作后  
两个进程同时  
在管程中

- Mutual exclusion within a monitor is ensured.



# Monitor Implementation

---

- For each condition variable  $x$ , we have

```
semaphore x-sem; // (initially = 0)
int x-count = 0;
```

挂起等待资源  
的进程

- The operation  $x.\text{wait}$  can be implemented as

```
x-count++;
if (next-count > 0)
 signal (next);
else
 signal (mutex);
wait (x-sem);
x-count--;
```





# Monitor Implementation

---

- The operation `x.signal` can be implemented as

```
if (x-count > 0)
{
 next-count++;
 signal (x-sem);
 wait (next);
 next-count- -;
}
```



# Monitor Implementation

---

- ***Conditional-wait*** construct: `x.wait(c)`
  - `c` – integer expression evaluated when the wait operation is executed.
  - value of `c` (*a priority number*) stored with the name of the process that is suspended.
  - when `x.signal()` is executed, process with smallest associated priority number is resumed next.



# Monitor Implementation

---

- **Check two conditions to establish correctness of system**
  - **User processes must always make their calls on the monitor in a correct sequence.**
  - **Must ensure that an uncooperative process does not ignore the mutual-exclusion gateway provided by the monitor, and try to access the shared resource directly, without using the access protocols.**



# Outline

---

- **Background**
- **The Critical-Section Problem**
- **Synchronization Hardware**
- **Semaphores**
- **Classical Problems of Synchronization**
- **Monitors**
- **Synchronization Examples**



# Solaris Synchronization

---

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing.
- Uses ***adaptive mutex*** for efficiency when protecting data from short code segments.
- Uses ***condition variables*** , ***semaphore***, and ***readers-writers locks*** when longer sections of code need access to data.
- Uses ***turnstile*** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock.



# Windows XP Synchronization

---

- Uses **interrupt masks** (中断屏蔽) to protect access to global resources on uniprocessor systems.
- Uses **spinlocks** on multiprocessor systems.
- Also provides **dispatcher objects** which may act as mutex and semaphores.
- Dispatcher objects may also provide *events*. An event acts much like a condition variable.



# Linux Synchronization

---

- **Linux**
  - **disables interrupts to implement short critical sections**
- **Linux provides**
  - **semaphores**
  - **spinlocks**



# Pthreads Synchronization

---

- Pthreads API is OS-independent
- It provides
  - mutex locks
  - condition variables
- Non-portable extensions include
  - read-write locks
  - spinlocks