



操作系统原理及应用

李 伟

xchlw@seu.edu.cn

计算机科学与工程学院、软件学院
江苏省网络与信息安全重点实验室



Chapter 9 Virtual Memory



Outline


- **Background**
- **Demand Paging**
- **Page Replacement**
- **Allocation of Frames**
- **Thrashing**
- **Memory-Mapped Files**
- **Allocating Kernel Memory**
- **Other Considerations**
- **Operating-System Examples**




Background

- **Virtual memory – separation of user logical memory from physical memory.**
 - **Only part of the program needs to be in memory for execution.**
 - **Logical address space can therefore be much larger than physical address space.**
 - **More programs can be run at the same time**
 - **Less I/O be needed to load or swap**

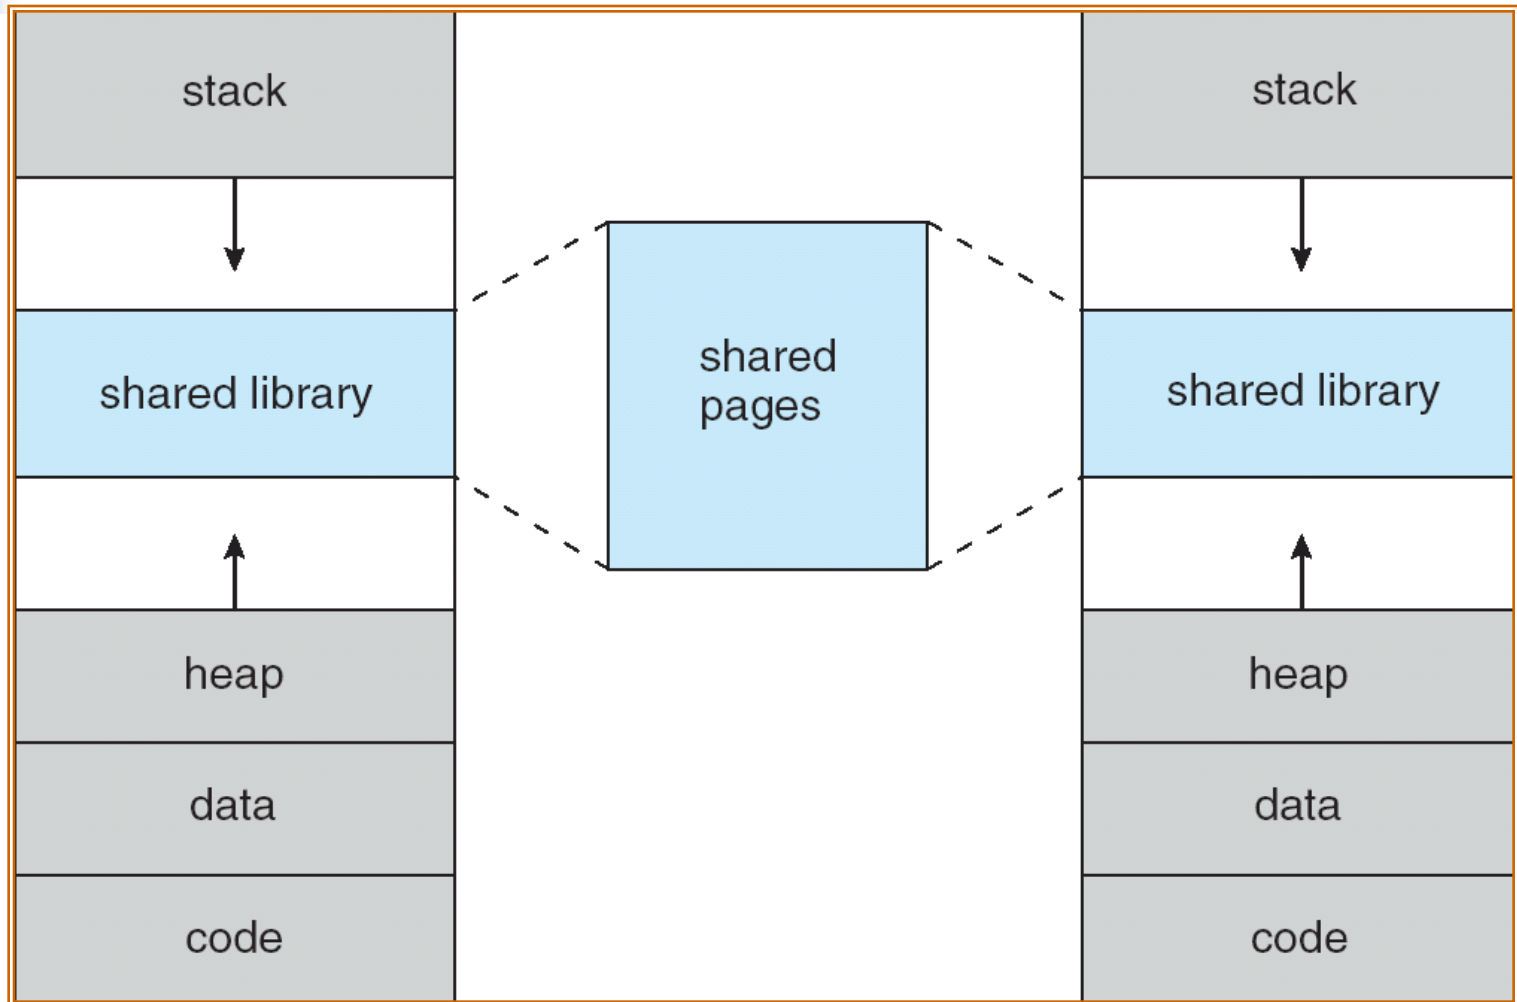
Virtual Memory is Larger than Physical Memory



Virtual-address Space



Shared Library Using Virtual Memory





Copy-on-Write

- Process Creation
- Copy-on-Write (COW) allows both parent and child processes to **initially share** the same pages in memory.
- If either process modifies a shared page, then only the page is copied.
- COW allows more efficient process creation as only modified pages are copied.
- Free pages are allocated from **a pool of zeroed-filled pages**.



Outline

- Background
- **Demand Paging**
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples



Demand Paging

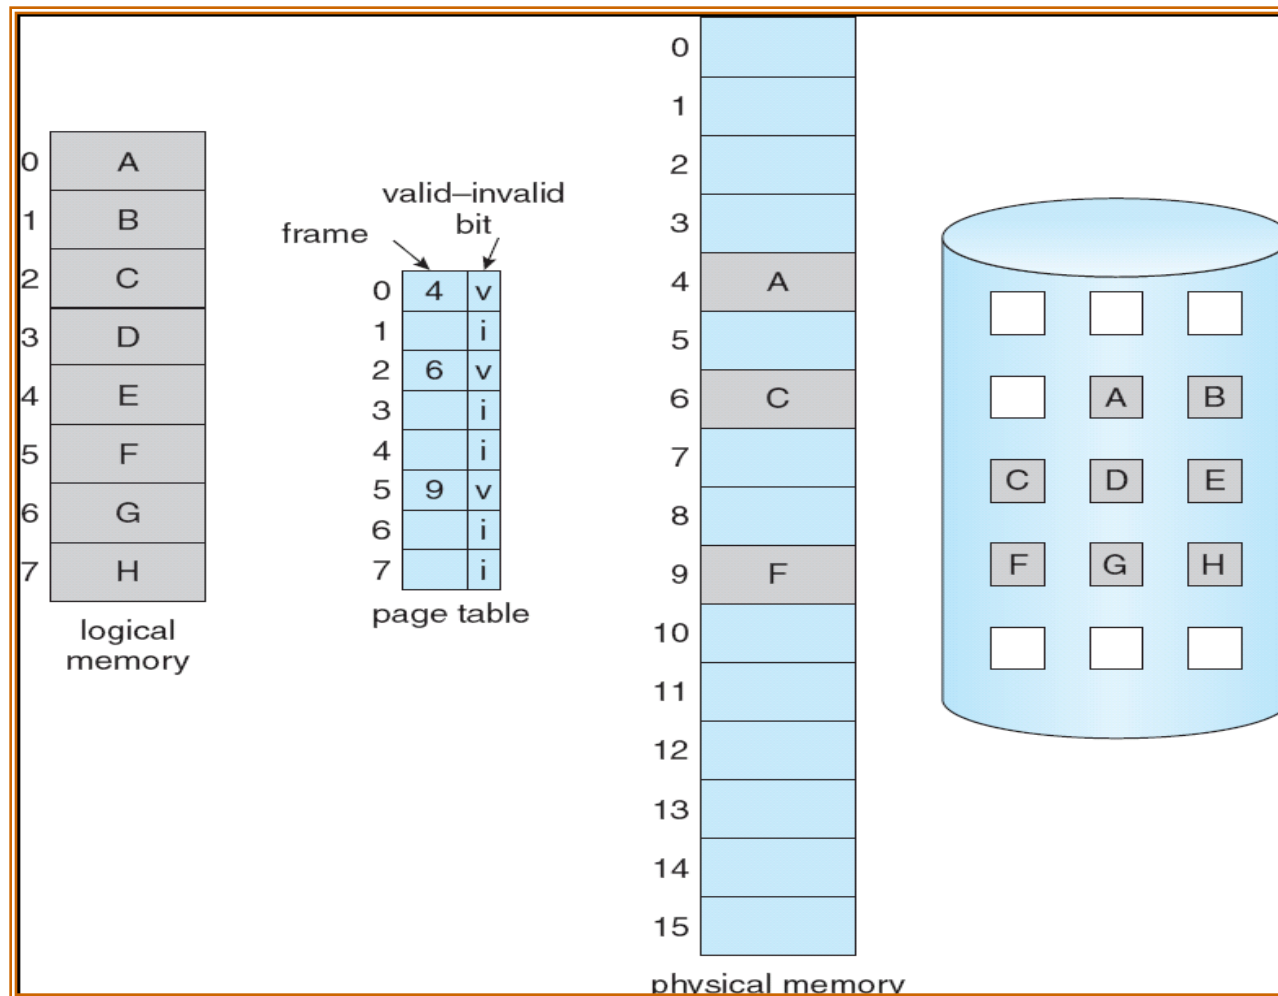
- **Bring a page into memory only when it is needed.**
 - **Less I/O needed**
 - **Less memory needed**
 - **Faster response**
 - **More users**
- **Page is needed \Rightarrow reference to it**
 - **invalid reference \Rightarrow abort**
 - **not-in-memory \Rightarrow bring to memory**
- **Lazy Swapper (Pager) – never swaps a page into memory unless page will be needed**



Valid-Invalid Bit

- A valid–invalid bit is associated with each page table entry
 - 1 \Rightarrow valid and in-memory
 - 0 \Rightarrow invalid or not-in-memory
- Initially valid–invalid bit is set to 0 on all entries.
- During address translation, if valid–invalid bit in page table entry is 0 \Rightarrow **page fault**.

Page Table When Some Pages Are Not in Main Memory

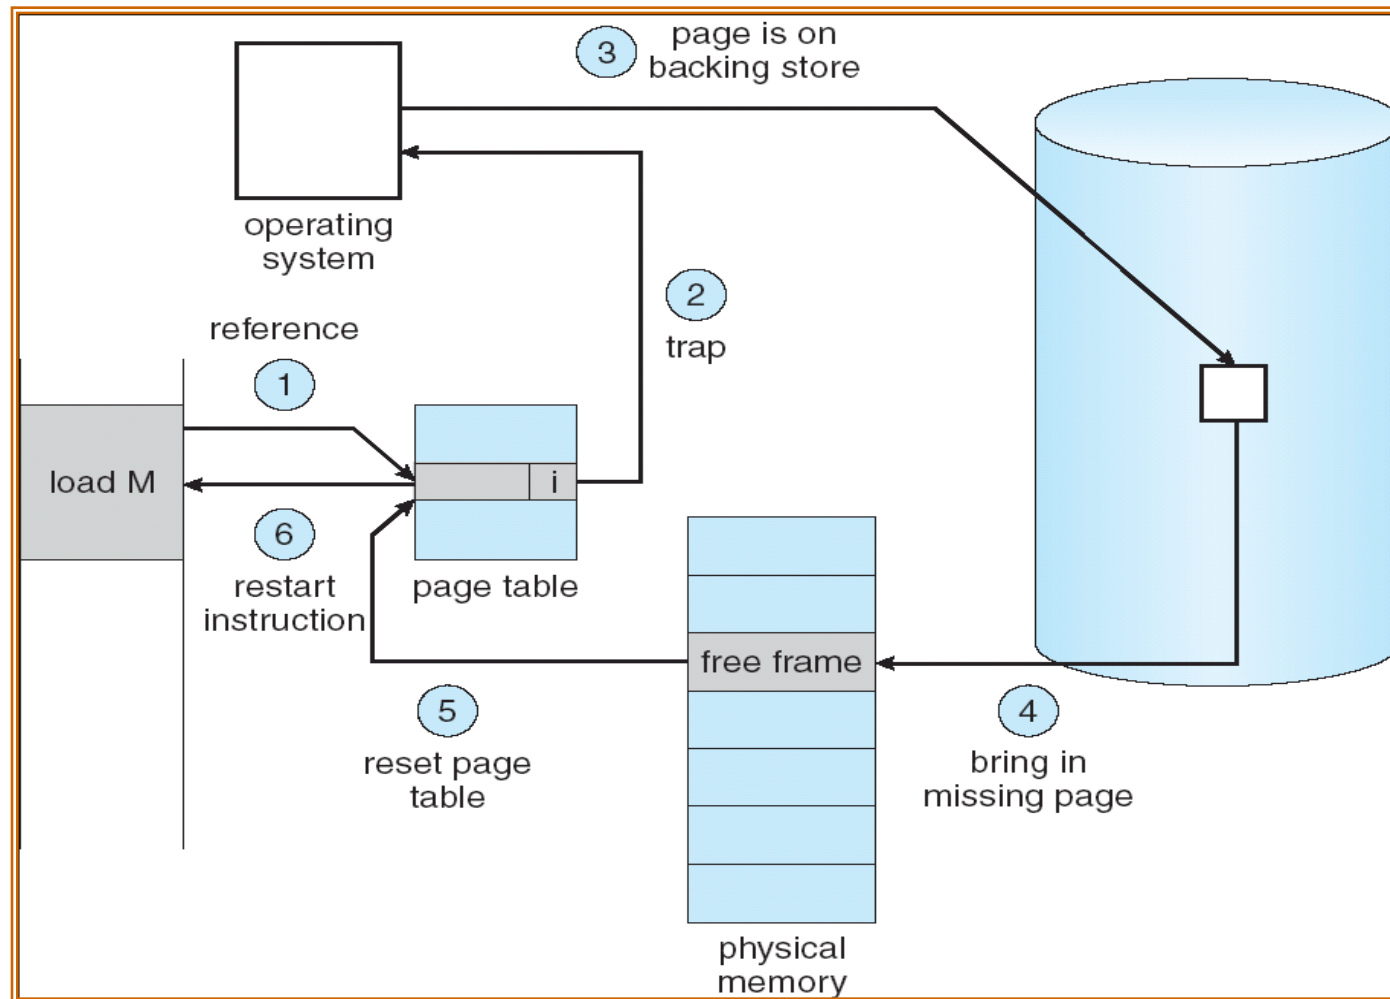




Page Fault

- If there is ever a reference to a page, first reference will trap to OS \Rightarrow page fault
 - OS looks at internal table to decide:
 - Invalid reference \Rightarrow abort.
 - Just not in memory.
 - Get empty frame.
 - Swap page into frame.
 - Reset tables, validation bit = 1.
 - Restart instruction

Steps in Handling a Page Fault



Performance of Demand Paging

- **Page Fault Rate** $0 \leq p \leq 1$

- if $p = 0$ no page faults
- if $p = 1$, every reference is a fault

- **Effective Access Time (EAT)**

$$\text{EAT} = (1 - p) \times \text{memory access time} \\ + p (\text{page fault overhead})$$

page fault overhead = service the page-fault interrupt
+ [swap page out]
+ swap page in
+ restart overhead



Outline

- Background
- Demand Paging
- **Page Replacement**
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples



What happens if there is no free frame?


- Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement.
- Page replacement – find some page in memory, but not really in use, swap it out
 - algorithm



Page Replacement

- Page replacement completes separation between logical memory and physical memory
 - large virtual memory can be provided on a smaller physical memory.
- Same page may be brought into memory several times
 - performance – want an algorithm which will result in minimum number of page faults

Need For Page Replacement






Basic Page Replacement

- Find the location of the desired page on disk.
- Find a free frame
 - If there is a free frame, use it.
 - If there is no free frame, use a **page replacement algorithm** to select a **victim frame**.
 - Write the victim frame to the disk and change the page and frame tables.
- Read the desired page into the free frame. Update the page and frame tables.
- Restart the process.

Basic Page Replacement




Use **modify (dirty) bit** to reduce overhead of page transfers
– only modified pages are written to disk.



Page Replacement Algorithms

- Want **lowest page-fault rate.**
- Evaluate algorithm by running it on a particular string of memory references (**reference string**) and computing the number of page faults on that string.

Graph of Page Faults Versus The Number of Frames



First-In-First-Out (FIFO) Algorithm

- Reference string
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

1	4	5
2	1	3
3	2	4

9 page faults


First-In-First-Out (FIFO) Algorithm

- 4 frames

1	5	4	
2	1	5	10 page faults
3	2		
4	3		

- FIFO Replacement – **Belady's Anomaly**
 - more frames \Rightarrow more page faults

FIFO Illustrating Belady's Anamoly





Optimal Algorithm

- Replace page that will not be used for longest period of time.
- 4 frames example — 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	4
2	
3	
4	5

6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs.

Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	5
2	
3	5 4
4	3

- **Counter implementation**
 - Every page entry has a **counter**; every time page is referenced through this entry, the clock is copied into the counter.
 - Replace the page with the smallest time value.




LRU Algorithm

- **Stack implementation** – keep a stack of page numbers in a double link form
 - Page referenced
 - remove it from the stack and put it on the top
 - LRU page is always at the bottom
 - requires 6 pointers to be changed at worst
 - No search for replacement
- Optimal replacement and LRU replacement do not suffer from Belady's anomaly

Use of a Stack to Record the Most Recent Page References

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b





LRU Approximation Algorithms

■ Reference bit

- With each entry in the page table associate a bit, initially = 0
- When page is referenced, the bit will be set to 1.
- Replace the one which bit is 0 (if one exists).
We do not know the order of use, however we can determine which pages have been used and which have not been used.



LRU Approximation Algorithms

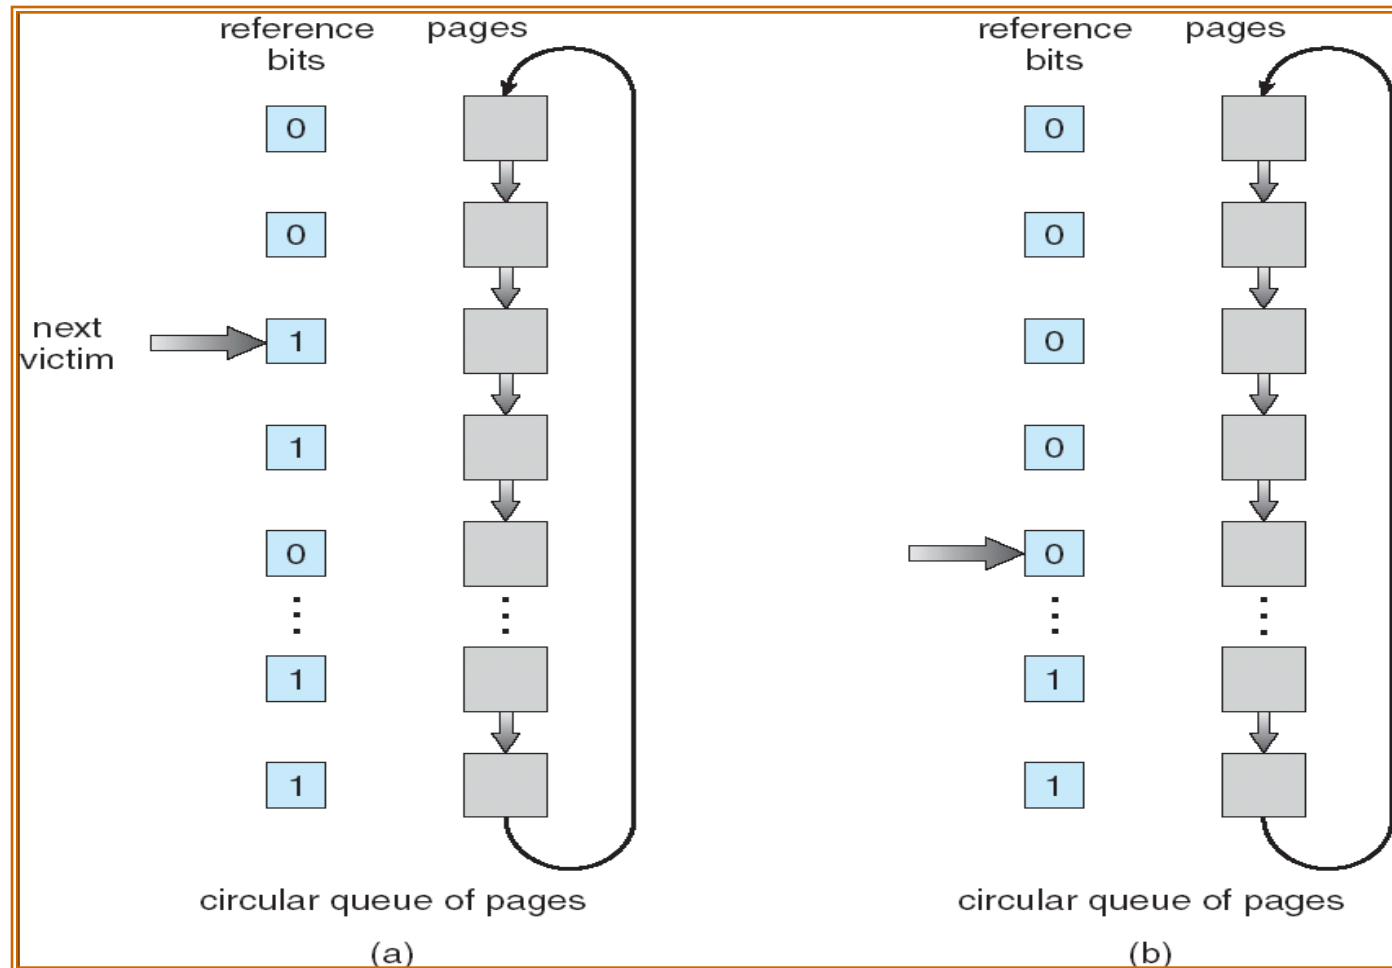
- **Additional-Reference-Bits Algorithm**
 - Keep an 8-bit bytes for each page
 - At regular intervals shifts the bits **right** 1 bit, shift the **reference bit** into the high-order bit
 - Interpret these 8-bit bytes as unsigned integers, **the page with lowest number** is the LRU page



LRU Approximation Algorithms

- **Second-Chance Algorithm**
 - Need reference bit.
 - Clock replacement (**FIFO**).
 - If page to be replaced (in clock order) has reference bit = 1. then:
 - set reference bit 0.
 - leave page in memory.
 - replace next page (in clock order), subject to same rules.

Second-Chance Page-Replacement Algorithm





Counting Algorithms

- Keep a counter of the number of references that have been made to each page.
- **LFU Algorithm:** replaces page with smallest count.
- **MFU Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used.



作业1

- 在某请求分页管理系统中，一个作业共**5**页，作业执行时依次访问如下页面：**1、4、3、1、2、5、1、4、2、1、4、5**，若分配给该作业的主存块数为**3**，分别采用**FIFO**、**LRU**、**OPT**页面置换算法，试求出缺页中断的次数及缺页率。



Outline

- Background
- Demand Paging
- Page Replacement
- **Allocation of Frames**
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples



Allocation of Frames

- Each process needs **minimum** number of pages, **which is decided by the given computer architecture.**
- Example: IBM 370 MVC instruction:
 - instruction is 6 bytes, might span 2 pages.
 - 2 pages to handle from.
 - 2 pages to handle to.
- Two major allocation schemes.
 - fixed allocation
 - priority allocation



Fixed Allocation

- **Equal allocation** – e.g., if 100 frames and 5 processes, give each 20 frames.
- **Proportional allocation** – Allocate according to the size of process.

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$



Priority Allocation

- Use a proportional allocation scheme using priorities rather than size.
- If process P_i generates a page fault,
 - select for replacement one of its frames.
 - select for replacement a frame from a process with lower priority number.



Global vs. Local Allocation

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- **Local replacement** – each process selects from only its own set of allocated frames.



Outline

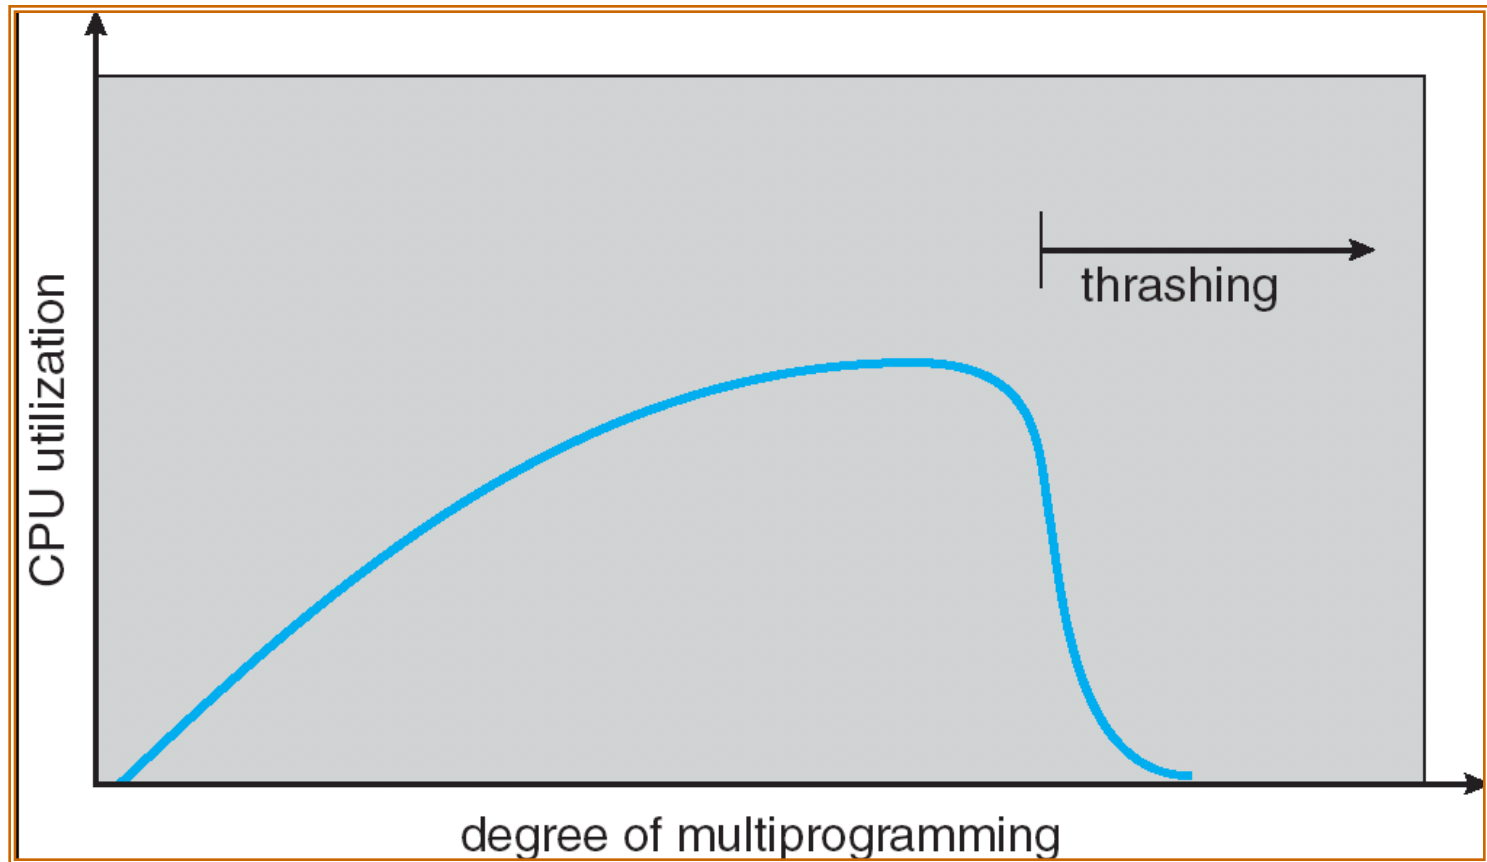
- Background
- Demand Paging
- Page Replacement
- Allocation of Frames
- **Thrashing**
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples



Thrashing (颠簸)

- If a process does not have “enough” frames, the page-fault rate is very high.
 - low CPU utilization.
 - operating system thinks that it needs to increase the degree of multiprogramming.
 - another process is added to the system.
- **Thrashing** —— a process is busy swapping pages in and out.

Thrashing





Thrashing

- **Why does paging work?**

- Locality model**

- A locality is a set of pages that are actively used together.
 - Process migrates from one locality to another.
 - Localities may overlap.

- **Why does thrashing occur?**

- size of locality > allocated memory size



Working-Set Model

- $\Delta \equiv$ **working-set window** \equiv a fixed number of page references
Example: 10,000 instruction
- WSS_i (**working set** of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass entire locality.
 - if Δ too large will encompass several localities.
 - if $\Delta = \infty \Rightarrow$ will encompass entire program.




Working-Set Model

- The working set is an approximation of the program's locality.
- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m \Rightarrow$ Thrashing \Rightarrow Suspend one of the processes


Working-set model

page reference table

... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...




$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

Page-Fault Frequency



- Establish “acceptable” page-fault rate.
 - If actual rate too low, process loses frame.
 - If actual rate too high, process gains frame.



Outline


- Background
- Demand Paging
- Page Replacement
- Allocation of Frames
- Thrashing
- **Memory-Mapped Files**
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples




Memory-Mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by *mapping* a disk block to a page in memory.
- A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical frame. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- Simplifies file access by treating file I/O through memory rather than read()/write() system calls.
- Also allows several processes to map the same file allowing the pages in memory to be shared.

Memory-Mapped Files



Memory-Mapped Shared Memory in Windows





Outline

- Background
- Demand Paging
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- **Allocating Kernel Memory**
- Other Considerations
- Operating-System Examples



Allocating Kernel Memory

- **Treated differently from user memory**
- **Often allocated from a free-memory pool**
 - **Kernel requests memory for structures of varying sizes**
 - **Some kernel memory needs to be contiguous**

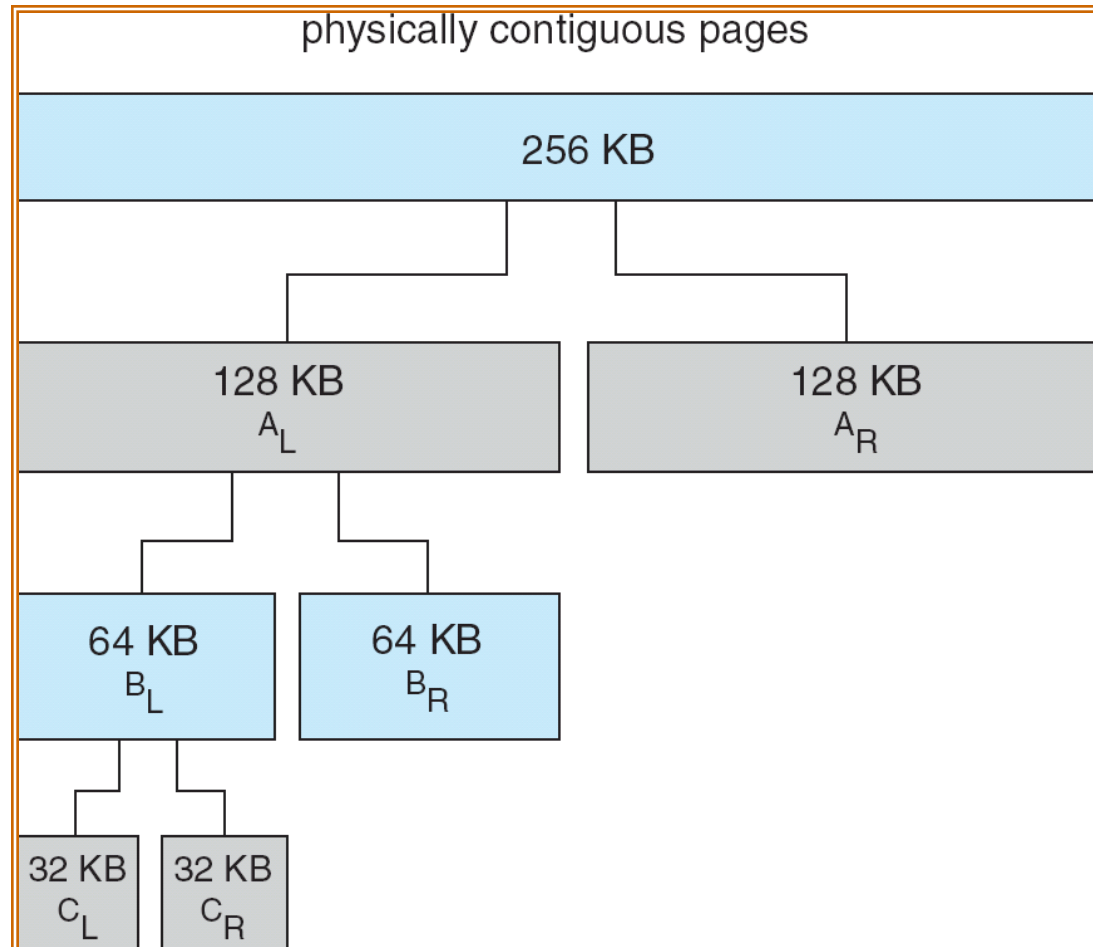


Buddy System

- **Allocates memory from fixed-size segment consisting of physically-contiguous frames**
- **Memory allocated using power-of-2 allocator**
 - **Satisfies requests in units sized as power of 2**
 - **Request rounded up to next-highest power of 2**
 - **When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2**
 - **Continue until appropriate sized chunk available**



Buddy System Allocator





Slab Allocator


- Alternate strategy
- **Slab** is one or more physically contiguous pages
- **Cache** consists of one or more slabs
- Single cache for each unique kernel data structure
 - Each cache filled with objects – instantiations of the data structure



Slab Allocator

- When cache created, filled with objects marked as free
- When structures stored, objects marked as used
- If slab is full of used objects, next object allocated from empty slab
 - If no empty slabs, new slab allocated
- Benefits include no fragmentation, fast memory request satisfaction

Slab Allocation





Outline

- **Background**
- **Demand Paging**
- **Page Replacement**
- **Allocation of Frames**
- **Thrashing**
- **Memory-Mapped Files**
- **Allocating Kernel Memory**
- **Other Considerations**
- **Operating-System Examples**



Other Issues -- Prepaging

- **Prepaging**
 - **To reduce the large number of page faults that occurs at process startup**
 - **Prepaging all or some of the pages a process will need, before they are referenced**
 - **But if prepaged pages are unused, I/O and memory was wasted**
 - **Assume s pages are prepaged and α percent of the pages is used**
 - **Is cost of $s * \alpha$ save pages faults $>$ or $<$ the cost of $s * (1 - \alpha)$ unnecessary pages?**
 - **α near 0 / 1 \Rightarrow prepaging loses / success**



Other Issues – Page Size

- **Page size selection must take into consideration**
 - **fragmentation (small page)**
 - **table size (large page)**
 - **I/O overhead (large page)**
 - **locality (small page)**



Other Issues – TLB Reach

- **TLB Reach - The amount of memory accessible from the TLB**
- **TLB Reach = (TLB Size) X (Page Size)**
- **Ideally, the working set of each process is stored in the TLB. Otherwise there is a high degree of page faults**
- **Increase the Page Size**
- **Provide Multiple Page Sizes**
 - **This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation**



Other Issues – Program Structure



- Program structure

- `int A[][] = new int[1024][1024];`

- Each row is stored in one page

- Program 1

```
for (j = 0; j < A.length; j++)  
    for (i = 0; i < A.length; i++)  
        A[i,j] = 0;
```

1024 x 1024 page faults

- Program 2

```
for (i = 0; i < A.length; i++)  
    for (j = 0; j < A.length; j++)  
        A[i,j] = 0;
```

- 1024 page faults



Other Issues – I/O interlock

- **I/O Interlock – Pages must sometimes be locked into memory**
- **Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm**



Outline

- Background
- Demand Paging
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- **Operating-System Examples**



Operating System Examples

- **Windows XP**
- **Solaris**



Windows XP

- **Uses demand paging with clustering.**
Clustering brings in pages surrounding the faulting page.
- **Processes are assigned working set minimum and working set maximum**
- **Working set minimum** is the minimum number of pages the process is guaranteed to have in memory



Windows XP

- A process may be assigned as many pages up to its working set maximum
- When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- Working set trimming removes pages from processes that have pages in excess of their working set minimum



Solaris


- Paging is performed by ***pageout* process**
- Pageout scans pages using **modified clock algorithm**
- ***Scanrate*** is the rate at which pages are scanned. This ranges from ***slowscan*** to ***fastscan***
- Pageout is called more frequently depending upon the amount of free memory available



Solaris

- Maintains **a list of free pages** to assign faulting processes
- *Lotsfree* – threshold parameter (amount of free memory) to begin paging
- *Desfree* – threshold parameter to increasing paging
- *Minfree* – threshold parameter to being swapping

Solaris 2 Page Scanner





Part 3 小结

- 基本概念：内存保护（基址寄存器+界限地址寄存器）、地址绑定、逻辑地址空间与物理地址空间、动态加载、动态链接、交换
- 连续分配：固定分区、可变分区（动态分配问题，3种方法）、碎片（内、外）
- 分页：页（页大小取值因素）、帧、页表、逻辑地址结构、页表实现（寄存器、PTBR、TLB）、Hit Ratio（命中率）、内存保护、共享页、页表结构（层次、哈希、反向）、分段
- 按需调页（有效访问时间）、写时复制、页置换算法（FIFO（Belady异常）、最优OPT、LRU、近似LRU、计数）、
- 帧分配（平均、按比例）、系统颠簸（含义、原因、解决方法）、内存映射文件、内核内存分配