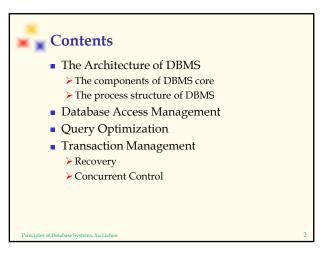
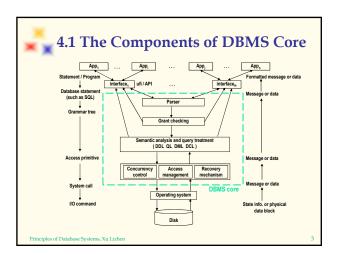
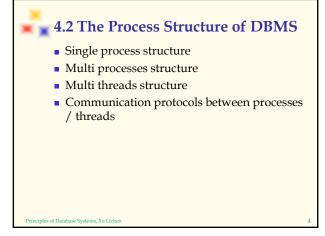
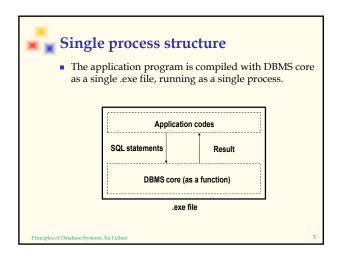
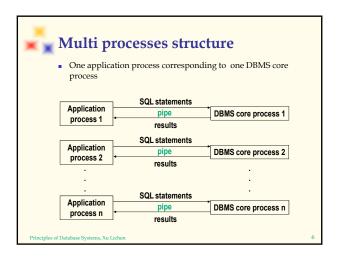
4. Database Management
Systems\*

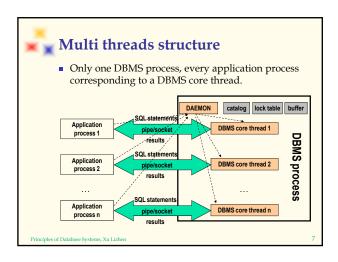














# 4.3 Database Access Management

The access to database is transferred to the operations on files (of OS) eventually. The file structure and access route offered on it will affect the speed of data access directly. It is impossible that one kind of file structure will be effective for all kinds of data access.

- Access types
- File organization
- Index technique
- Access primitives

Principles of Database Systems, Xu Lizher



# Access Types

- Query all or most records of a file (>15%)
- Query some special record
- Query some records (<15%)</li>
- Scope query
- Update

Principles of Database Systems, Xu Lizhe



# File Organization

- Heap file: records stored according to their inserted order, and retrieved sequentially. This is the most basic and general form of file organization.
- Direct file: the record address is mapped through hash function according to some attribute's value.
- Indexed file: index + heap file/cluster
- Dynamic hashing: p115
- Grid structure file: p118 (suitable for multi attributes queries)
- Raw disk (notice the difference between the logical block and physical block of file. You can control physical blocks in OS by using raw disk)

Principles of Database Systems, Xu Lizher

10



### Index Technique

- B+ Tree (√√)
- Clustering index ( √)
- Inverted file
- Dynamic hashing
- Grid structure file and partitioned hash function
- Bitmap index (used in data warehouse)
- Others

\_\_\_\_\_\_



# 4.4 Query Optimization

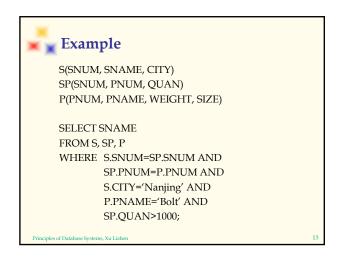
"Rewrite" the query statements submitted by user first, and then decide the most effective operating method and steps to get the result. The goal is to gain the result of user's query with the lowest cost and in shortest time.

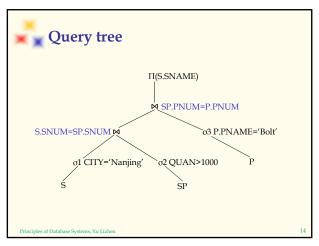
#### 4.4.1 Summary of Query Optimization

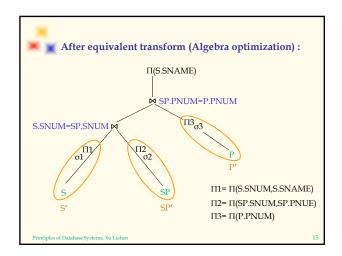
- Algebra Optimization
- Operation Optimization

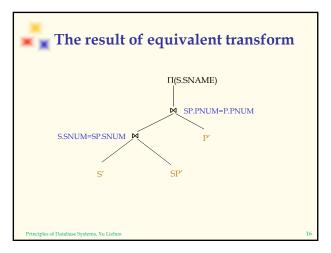
Principles of Database Systems, Xu Lizhen

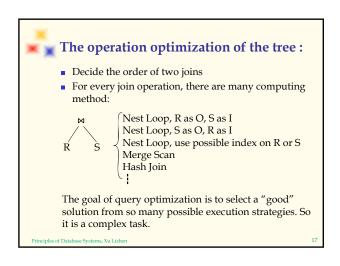
12

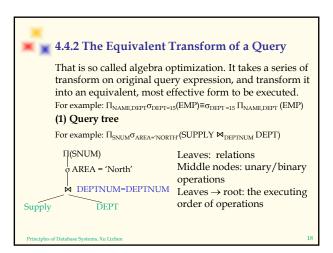


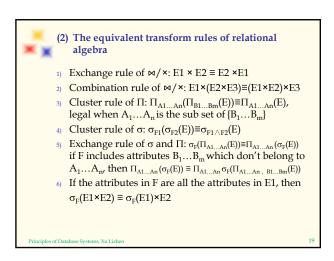


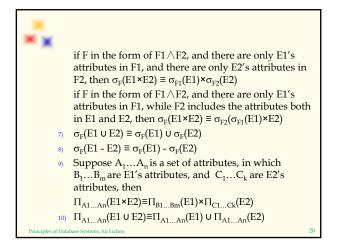




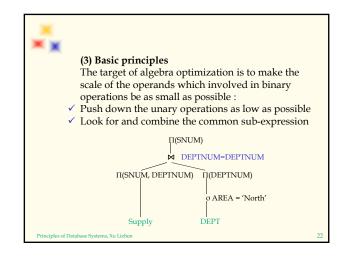








SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT \*
FROM Reserves R
WHERE R.bid=103 AND S.sid=R.sid)



# 4.4.3 The Operation Optimization How to find a "good" access strategy to compute the query improved by algebra optimization is introduced in this section: Optimization of select operation Optimization of project operation Optimization of set operation Optimization of join operation Optimization of combined operations

Optimization of join operation

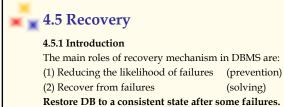
Nested loop: one relation acts as outer loop relation (O), the other acts as inner loop relation (I). For every tuple in O, scan I one time to check join condition.

Because the relation is accessed from disk in the unit of block, we can use block buffer to improve efficiency. For R ⋈ S, if let R as O, S as I, b<sub>R</sub> is physical block number of R, b<sub>S</sub> is physical block number of S, there are n<sub>B</sub> block buffers in system (n<sub>B</sub>>=2), and n<sub>B</sub>-1 buffers used for O, one buffer used for I, then the total disk access times needed to compute R ⋈ S is:
b<sub>R</sub> + rb<sub>R</sub>/(n<sub>B</sub>-1)1×b<sub>S</sub>



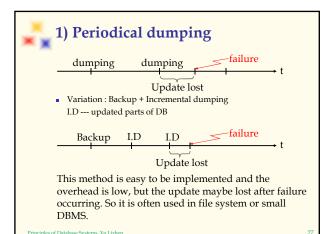
# Optimization of join operation

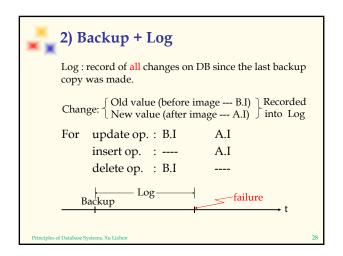
- Merge scan: order the relation R and S on disk in ahead, then we can compare their tuples in order, and both relation only need to scan one time. If R and S have not ordered in ahead, must consider the ordering cost to see if it is worth to use this method
- Using index or hash to look for mapping tuples: in nested loop method, if there is suitable access route on I (say B+ tree index), it can be used to substitute sequence scan. It is best when there is cluster index or hash on join attributes.
- Hash join: because the join attributes of R and S have the same domain, R and S can be hashed into the same hash file using the same hash function, then R ⋈ S can be computed based on the hash file.



- Redundancy is necessary.
- Should inspect all possible failures.
- General method:

(solving)







#### While recovering:

- Some transactions maybe half done, should undo them with B.I recorded in Log.
- Some transactions have finished but the results have not been written into DB in time, should redo them with A.I recorded in Log. (finish writing into DB)

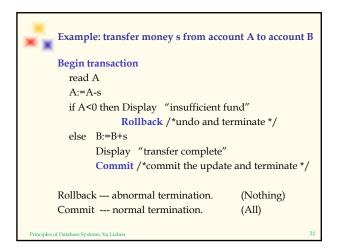
It is possible to recover DB to the most recent consistent state with Log.

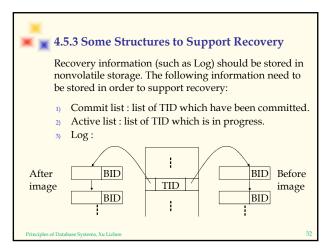


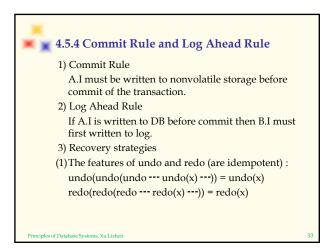
#### 4.5.2 Transaction

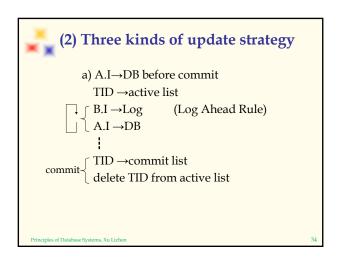
A transaction T is a finite sequence of actions on DB exhibiting the following effects:

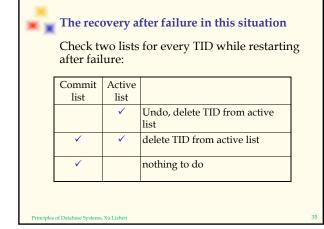
- Atomic action: Nothing or All.
- Consistency preservation: consistency state of  $DB \rightarrow$  another consistency state of DB.
- Isolation: concurrent transactions should run as if they are independent each other.
- Durability: The effects of a successfully completed transaction are permanently reflected in DB and recoverable even failure occurs later.

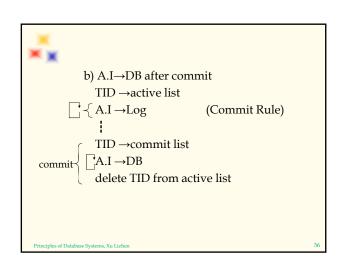


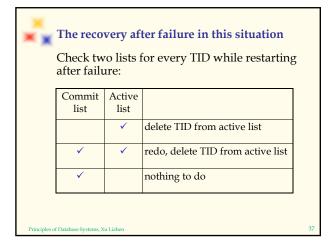


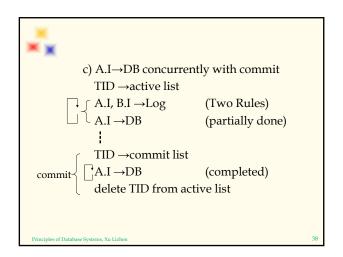


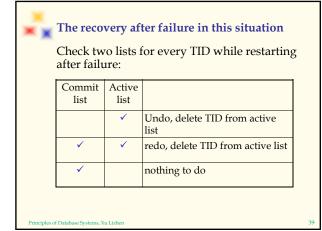


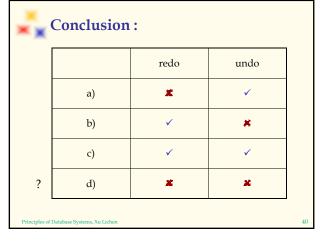


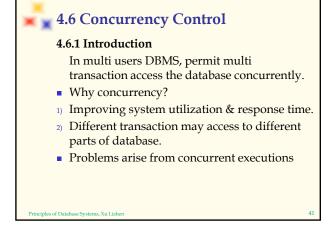


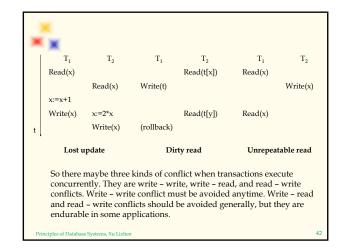










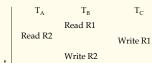




# 4.6.2 Serialization --- the criterion for concurrency consistency

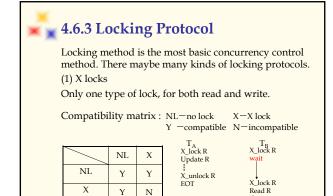
Definition: suppose  $\{T_1,T_2,...T_n\}$  is a set of transactions executing concurrently. If a schedule of  $\{T_1,T_2,...T_n\}$  produces the same effect on database as some serial execution of this set of transactions, then the schedule is serializable.

Problem: different schedule  $\rightarrow$  different equivalent serial execution  $\rightarrow$  different result? (yes, n!)



The result of this schedule is the same as serial execution  $T_A \rightarrow T_B \rightarrow T_{C'}$  so it is serializable. The equivalent serial execution is  $T_A \rightarrow T_B \rightarrow T_C$ .

Principles of Database Systems, Xu Lizhen





## \*Two Phase Locking

- Definiton1: In a transaction, if all locks precede all unlocks, then the transaction is called two phase transaction. This restriction is called two phase locking protocol.
- Definition2: In a transaction, if it first acquires a lock on the object before operating on it, it is called wellformed.
- Theorem: If S is any schedule of well-formed and two phase transactions, then S is serializable. (proving is on p151)

Growing by Lock A Lock B Lock C Unlock A Shrinking Unlock B Unlock C 2PI

T<sub>2</sub>
Lock A
Lock B
Unlock A
Unlock B
Lock C
Unlock C

Principles of Database Systems, Xu Lizher



- 1) Well-formed + 2PL: serializable
- Well-formed + 2PL + unlock update at EOT: serializable and recoverable. (without domino phenomena)
- Well-formed + 2PL + holding all locks to EOT: strict two phase locking transaction.
- (2) (S,X) locks

S lock --- if read access is intended.

X lock --- if update access is intended.

	NL	S	Χ
NL	Y	Y	Y
S	Y	Y	N
X	Y	N	N

Principles of Database Systems, Xu Lizher



# (3) (S,U,X) locks

U lock --- update lock. For an update access the transaction first acquires a U-lock and then promote it to X-lock. Purpose: shorten the time of exclusion, so as to boost concurrency degree, and reduce deadlock.

	NL	S	U	Χ
NL	Y	Y	Y	Y
S	Y	Y	Y	N
U	Y	Y	N	N
X	Y	N	N	N

X (S,X) (S,U,X)

Concurrency degree

Overhead

Principles of Database Systems, Xu Lizhe



## 4.6.4 Deadlock & Live Lock

Dead lock: wait in cycle, no transaction can obtain all of resources needed to complete.

Live lock: although other transactions release their resource in limited time, some transaction can not get the resources needed for a very long time.



- Live lock is simpler, only need to adjust schedule strategy, such as FIFO
- Deadlock: (1) Prevention(don't let it occur); (2) Solving(permit it occurs, but can solve it)

rinciples of Database Systems, Xu Lizhen



# (1) Deadlock Detection

- Timeout: If a transaction waits for some specified time then deadlock is assumed and the transaction should be aborted.
- 2) Detect deadlock by wait-for graph G= $\langle V,E \rangle$ V: set of transactions  $\{T_i \mid T_i \text{ is a transaction in DBS } (i=1,2,...n)\}$ 
  - $E : \{ \langle T_i, T_j \rangle \mid T_i \text{ waits for } T_j \text{ } (i \neq j) \}$
- If there is cycle in the graph, the deadlock occurs.
- When to detect?
- whenever one transaction waits.
- 2) periodically

rinciples of Database Systems, Xu Lizhen



- What to do when detected?
- 1) Pick a victim (youngest, minimum abort cost, ...)
- 2) Abort the victim and release its locks and resources
- 3) Grant a waiter
- 4) Restart the victim (automatically or manually)
- (2) Deadlock avoidance
- 1) Requesting all locks at initial time of transaction.
- 2) Requesting locks in a specified order of resource.
- 3) Abort once conflicted.
- 4) Transaction Retry

Principles of Database Systems, Xu Lizhen

50



Every transaction is uniquely time stamped. If  $T_A$  requires a lock on a data object that is already locked by  $T_{Br}$  one of the following methods is used:

- a) Wait-die:  $T_A$  waits if it is older than  $T_B$ , otherwise it "dies", i.e. it is aborted and automatically retried with original timestamp.
- b) Wound-wait:  $T_A$  waits if it is younger than  $T_B$ , otherwise it "wound"  $T_B$ , i.e.  $T_B$  is aborted and automatically retried with original timestamp.

In above, both have only one direction wait, either older  $\rightarrow$  younger or younger  $\rightarrow$  older. It is impossible to occur wait in cycle, so the dead lock is avoided.

Principles of Database Systems, Xu Lizhe

51